

---

# Imposing a Memory Management Discipline on Software Deployment

Eelco Dolstra, Eelco Visser and Merijn de Jonge  
Utrecht University, P.O. Box 80089,  
3508 TB Utrecht, The Netherlands  
{eelco, visser, mdejonge}@cs.uu.nl

March 22, 2004

---

## The problem


---

Software deployment (the act of transferring software to another system) is surprisingly hard.

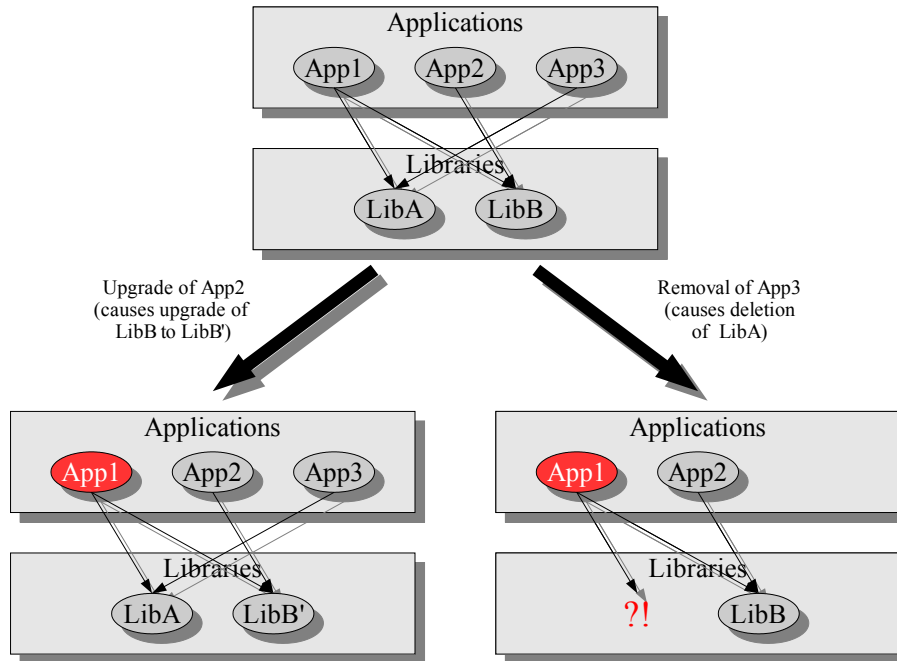
- Must ensure correctness.
  - Dependency information must be complete.
  - Component compatibility.
  - Atomicity of upgrades/downgrades.
  - Safe removal of unused components.

- 
- Lot of effort.
    - Packaging is often (semi-)manual.
    - Source/binary distributions.
    - Must package each variant.
    - Don't want to install all component separately.
    - Especially a problem with small-grained reuse (e.g., StrategoXT).
  - Should support multiple versions/variants.
    - Test a component before production use.
    - Multiple users.

# Incomplete Dependencies



# Interference



---

## The core problems

---

- Must prevent *unresolved component dependencies*.
  - A component should never refer to another component not present on the target system.
  - Hard to validate; how to detect use of undeclared dependencies?
  - Timeline issues: (related) dependencies at build and run time.
- Must prevent *component interference*.
  - Different versions/variants of a component (or completely unrelated components) should not interfere with each other.
  - Upgrades are usually *destructive*. E.g., only one `/usr/bin/gcc`.

---

# Software deployment as a memory-management problem

---

|                           |   |                               |
|---------------------------|---|-------------------------------|
| memory                    | ⇔ | disk                          |
| objects (values)          | ⇔ | components                    |
| addresses                 | ⇔ | path names                    |
| pointers are numbers      | ⇔ | pointers are strings          |
| pointer dereference       | ⇔ | I/O                           |
| pointer arithmetic        | ⇔ | string operations             |
| dangling pointer          | ⇔ | reference to absent component |
| object graph              | ⇔ | dependency graph              |
| persistence/serialisation | ⇔ | deployment                    |

---

## Closures

---

- Correct deployment of component  $c$  requires distributing the smallest set of components  $C$  containing  $c$  closed under the “has-a-pointer-to” relation.
- I.e., we have to discover the pointer graph.



---

## Determining the pointer graph

---

- This is just what garbage collectors for programming languages have to do.
- GC requires a *pointer discipline*:
  - Ideally, entire memory layout is known, and no arbitrary pointer formation (e.g., integer  $\Leftrightarrow$  pointer casts).
  - But even C/C++ has rules: pointer arithmetic is not allowed to move a pointer out of the object it points to.
  - This is why *conservative GC* works: assume that everything that looks like a pointer *is* a pointer.

- 
- However, software components do not have any pointer discipline.
    - Any string can be a pointer.
    - Pointer arithmetic and dereferencing directories can produce pointers to any object in the file system.

---

## A pointer discipline

---

Solution: *impose* a pointer discipline.

- Each component should include in its a path a unique identifying string.

`/nix/store/15373f8c93776a3a5f86fec65914e59d-subversion-0.37.0`

`/nix/store/b70b48128d8d13725346684ea43963c4-strategoxt-0.9.3`

- Then we can apply conservative GC techniques to determine the pointer graph.

---

## Scanning for pointers

---

```
080  00 80 04 08 34 41 01 00 34 41 01 00 05 00 00 00 |....4A..4A.....|
090  00 10 00 00 01 00 00 00 34 41 01 00 34 d1 05 08 |.....4A..4...|
0a0  34 d1 05 08 b4 04 00 00 c4 04 00 00 06 00 00 00 |4.....|
0b0  00 10 00 00 02 00 00 00 7c 41 01 00 7c d1 05 08 |.....|A..|...|
0c0  7c d1 05 08 90 01 00 00 90 01 00 00 06 00 00 00 ||.....|
0d0  04 00 00 00 04 00 00 00 60 01 00 00 60 81 04 08 |.....'...'...|
0e0  60 81 04 08 20 00 00 00 20 00 00 00 04 00 00 00 |'.....|
0e0  60 81 04 08 20 00 00 00 20 00 00 00 04 00 00 00 |'.....|
0f0  04 00 00 00 50 e5 74 64 20 41 01 00 20 c1 05 08 |....P.td A.....|
100  20 c1 05 08 14 00 00 00 14 00 00 00 04 00 00 00 |.....|
110  04 00 00 00 2f 6e 69 78 2f 73 74 6f 72 65 2f 38 |..../nix/store/8|
120  64 30 31 33 65 61 38 37 38 64 30 66 66 38 34 63 |d013ea878d0ff84c|
130  62 31 37 38 61 34 62 31 36 30 65 34 30 32 36 2d |b178a4b160e4026-|
140  67 6c 69 62 63 2d 32 2e 33 2e 32 2f 6c 69 62 2f |glibc-2.3.2/lib/|
150  6c 64 2d 6c 69 6e 75 78 2e 73 6f 2e 32 00 00 00 |ld-linux.so.2...|
160  04 00 00 00 10 00 00 00 01 00 00 00 47 4e 55 00 |.....GNU.|
170  00 00 00 00 02 00 00 00 02 00 00 00 05 00 00 00 |.....|
180  83 00 00 00 bb 00 00 00 58 00 00 00 ab 00 00 00 |.....X.....|
190  ae 00 00 00 a1 00 00 00 00 00 00 00 6c 00 00 00 |.....1...|
```

---

## Risks

---

- Like all conservative GC approaches, there is a risk of *pointer hiding*.
  - Compressed executables.
  - UTF-16 encoded paths.
- Hasn't happened yet, though.

---

## Persistence

---

- The unique strings should be cryptographic hashes of all inputs involved in building the component.
- This prevents address collisions in the target address space (i.e., path name collisions in the target file system).

---

## Nix expressions

---

Component description in a pure functional language.

```
{stdenv, fetchurl, aterm, sdf}:

derivation {
  name = "strategoxt-0.9.3";
  system = stdenv.system;
  builder = ./builder.sh;
  src = fetchurl {
    url = ftp://.../strategoxt-0.9.3.tar.gz;
    md5 = "3425e7ae896426481bd258817737e3d6";
  };
  inherit stdenv, aterm, sdf;
}
```

---

## Nix expressions (2)

---

Build script:

```
#!/bin/sh

buildinputs="$aterm $sdf"
. $stdenv/setup || exit 1

tar zxf $src || exit 1
cd stratego* || exit 1
./configure --prefix=$out --with-aterm=$aterm \
  --with-sdf=$sdf || exit 1
make || exit 1
make install || exit 1
```



---

## Nix expressions (3)

---

### Composition: (all-packages.nix)

```
rec {
  strategox = (import ../development/compilers/strategox) {
    inherit fetchurl stdenv aterm;
    sdf = sdf2;
  };
  aterm = (import ../development/libraries/aterm) {
    inherit fetchurl stdenv;
  };
  sdf2 = (import ../development/tools/parsing/sdf2) {
    inherit fetchurl stdenv aterm getopt;
  };
  stdenv = ...;
  ...
}
```

---

## User operations

---

To build and install StrategoXT:

```
$ nix-env -if .../all-packages.nix strategoxt
```

When a new version comes along:

```
$ nix-env -uf .../all-packages.nix strategoxt
```

If it doesn't work:

```
$ nix-env --rollback
```

Delete unused components:

```
$ nix-collect-garbage
```

---

## Transparent binary deployment

---

On the producer side:

```
$ nix-push $(nix-instantiate .../all-packages.nix) \  
    http://server/cache
```

On the client side:

```
$ nix-pull http://server/cache
```

Installation will now reuse pre-built components, *iff* they are exactly the same.

---

## Implementation

---

- All components are stored in a *store* (e.g., `/nix/store`).
- Creation of components within the store described using *store expressions*.
- Store expressions describe a component build (a *derivation*) or the result thereof (a *closure*).
- Nix expressions are translated into store expressions. The path of the component is a cryptographic hash of *all* inputs into the build process. This ensures that no collisions occur between components.

---

## Conclusion

---

- Concurrent installation of multiple versions and variants.
- Atomic upgrades and downgrades.
- Multiple user environments.
- Safe dependencies.
- Complete deployment.
- Transparent source and binary deployment.
- Safe garbage collection.
- Portability.

---

## More information

---

- Website: <http://www.cs.uu.nl/groups/ST/Trace/Nix>.
- Eelco Dolstra, Eelco Visser and Merijn de Jonge. *Imposing a Memory Management Discipline on Software Deployment*. In *26th International Conference on Software Engineering (ICSE-2004)*, May 2004, Edinburgh (to appear).