

Purely Functional System Configuration Management

Software Technology Colloquium, UU

Eelco Dolstra¹ Armijn Hemel²

¹Universiteit Utrecht, Faculty of Science,
Department of Information and Computing Sciences

²Loohuis Consulting

April 19, 2007




Universiteit Utrecht

- ▶ Operating systems are installed and managed using tools that have an *imperative model*
- ▶ This causes lots of problems: upgrading is unreliable, rollbacks are hard, etc.
- ▶ This paper shows that it is possible to implement a system with a *purely functional model*
- ▶ Implemented in a Linux distribution called *NixOS*

What is a system configuration?

- ▶ Software packages
- ▶ Configuration files
- ▶ System scripts
- ▶ ...

Example of a configuration



Imperative model

System configuration tools have an *imperative model*: configuration actions are *stateful*.

Examples

- ▶ Package management tools such as RPM, apt, Windows installers perform *destructive updates*: they overwrite existing files.
- ▶ Configuration tools such as Cfengine (declaratively) specify imperative updates to configuration files.
- ▶ Windows installers overwrite registry entries.

Why is statefulness bad?

- ▶ No *traceability*
 - ▶ Configuration is the result of a sequence of (sometimes manual) imperative actions over time
 - ▶ Hard to reproduce a configuration
- ▶ No *predictability* (determinism)
 - ▶ If an action depends on an ill-defined initial state, then the result is probably ill-defined
 - ▶ This is why upgrading is riskier than a full re-install
- ▶ Configuration actions clobber the previous configuration
 - ▶ No rollbacks
 - ▶ Hard to safely test a configuration

Analogous to imperative languages

Imperative languages such as C and Java have analogous problems. E.g., cannot reason about the result of function calls due to global variables or I/O.

There is a better way!

- ▶ *Purely functional languages* like Haskell
 - ▶ No mutable variables, data structures
 - ▶ Function result only depends on function arguments
 - ▶ $x = y \Rightarrow f(x) = f(y)$
 - ▶ *Referential transparency*
- ▶ No referential transparency in existing system CM tools
- ▶ So we need purely function system configuration management!

There is a better way!

Goal: purely functional system configuration management

- ▶ All static parts of a configuration should be *immutable*
- ▶ Configurations should be built by *pure functions*

Nix: Purely functional package management

- ▶ Deployment system developed at Utrecht University:
<http://nix.cs.uu.nl/>
- ▶ *Purely functional* package management:
 - ▶ Packages builds only depend on declared inputs
 - ▶ Packages never change after they have been built

Nix store

Central idea: store all packages
in isolation from each other:

`/nix/store/axrzx0rh0ivw...
-firefox-2.0.0.3`

Paths contain a 160-bit
cryptographic hash of **all**
inputs used to build the
package:

- ▶ Sources
- ▶ Libraries
- ▶ Compilers
- ▶ Build scripts
- ▶ ...

```
/nix/store
├── l9w6773m1msy...-openssh-4.6p1
│   ├── bin
│   │   └── ssh
│   └── sbin
│       └── sshd
├── smkabrbibqv7...-openssl-0.9.8e
│   └── lib
│       └── libssl.so.0.9.8
├── c6jbqm2mc0a7...-zlib-1.2.3
│   └── lib
│       └── libz.so.1.2.3
└── im276akmsrhv...-glibc-2.5
    ├── lib
    │   └── libc.so.6
```

Nix expressions

Nix expressions describe how to build packages.

```
{stdenv, fetchurl, openssl, zlib}:
```

```
stdenv.mkDerivation {  
  name = "openssh-4.6p1";  
  src = fetchurl {  
    url = http://.../openssh-4.6p1.tar.gz;  
    sha256 = "0fpjlr3bfind0y94bk442x2p...";  
  };  
  buildCommand = "  
    tar xjf $src  
    ./configure --prefix=$out --with-openssl=${openssl}  
    make; make install ";  
}
```

Nix expressions

Nix expressions describe how to build packages.

```
{stdenv, fetchurl, openssl, zlib}:
```

Function arguments

```
stdenv.mkDerivation {  
  name = "openssh-4.6p1";  
  src = fetchurl {  
    url = http://.../openssh-4.6p1.tar.gz;  
    sha256 = "0fpjlr3bfind0y94bk442x2p...";  
  };  
  buildCommand = "  
    tar xjf $src  
    ./configure --prefix=$out --with-openssl=${openssl}  
    make; make install ";  
}
```

Nix expressions

Nix expressions describe how to build packages.

```
{stdenv, fetchurl, openssl, zlib}:
```

Function arguments

```
stdenv.mkDerivation {  
  name = "openssh-4.6p1";  
  src = fetchurl {  
    url = http://.../openssh-  
    sha256 = "0fpjlr3bfind0y94bk44zxzp...";  
  };  
  buildCommand = "  
    tar xjf $src  
    ./configure --prefix=$out --with-openssl=${openssl}  
    make; make install ";  
}
```

Build attributes

Nix expressions

system/all-packages.nix

```
openssh = import ../tools/networking/openssh {  
    inherit fetchurl stdenv openssl zlib;  
};  
  
openssl = import ../development/libraries/openssl {  
    inherit fetchurl stdenv;  
};  
  
zlib = import ../development/libraries/zlib {  
    inherit fetchurl stdenv;  
};  
  
fetchurl = ...;  
  
stdenv = ...;
```

Nix expressions

system/all-packages.nix

```
openssh = import ../tools/networking/openssh {  
    inherit fetchurl stdenv openssl zlib;  
};  
  
openssl = import ../development/libraries/openssl {  
    inherit fetchurl stdenv;  
};  
  
zlib = import ../development/libraries/zlib {  
    inherit fetchurl stdenv;  
};  
  
fetchurl = ...;  
  
stdenv = ...;
```



Taking it all the way

- ▶ Since we can build packages...
- ▶ ...why not build all the other stuff that goes into a configuration?
- ▶ It's all the same, really. As long as it's pure, we can build it!

The result: NixOS

- ▶ Linux distribution: didn't write anything ourselves, except for Nix and the necessary glue
- ▶ What we have:
 - ▶ Hardware support: networking, sound, video
 - ▶ RAID, LVM
 - ▶ System daemons: SSH, Apache, CUPS, dhcpd, NTP, Cron, Mingetty, ...
 - ▶ X11
 - ▶ KDE, most of Gnome
 - ▶ All the applications and tools in Nixpkgs (850 or so Unix packages)



What Nix expressions did we need?

- ▶ NixOS is currently 49 Nix expressions, about 6200 lines.
- ▶ Building Upstart jobs for starting system daemons / services.
- ▶ Building /etc configuration files.
- ▶ Building the *initial ramdisk* for booting.
- ▶ Building the Grub boot menu.
- ▶ Building the *activation script*.
- ▶ Building additional boot scripts.
- ▶ Building the ISO image (purely functional!).
- ▶ Building system management scripts: *nixos-rebuild*.

Example

/nix/store

└─ 068q49fhc600...-sshd

└─ etc/event.d

└─ sshd (*Upstart job*)

└─ lahm12vmh052...-sshd_config

└─ r1k7gb1capq0...-xauth-1.0.2

└─ bin

└─ xauth

└─ im276akmsrhv...-glibc-2.5


└─ lib

└─ libc.so.6

└─ l9w6773m1msy...-openssh-4.6p1

└─ sbin

└─ sshd



Example

/nix/store

└─ 068q49fhc600...-sshd

└─ etc/event.d

└─ sshd (*Upstart job*)

└─ lahm12vmh052...-sshd_config

└─ r1k7gb1capq0...-xauth-1.0.2

└─ bin

└─ xauth

lahm12vmh052...-sshd_config

UsePAM yes

X11Forwarding yes

XAuthLocation /nix/store/r1k7...-xauth-1.0.2/bin/xauth

...

└─ sshd

Example

Nix expression for sshd_config

```
{writeText, forwardX11, xauth}:
```

```
writeText "sshd_config" "
```

```
  UsePAM yes
```

```
  ${if forwardX11 then "
```

```
    X11Forwarding yes
```

```
    XAuthLocation ${xauth}/bin/xauth
```

```
  " else "
```

```
    X11Forwarding no
```

```
  "}
```

```
"
```

Example

Nix expression for the sshd Upstart job

```
{makeJob, openssh, sshdConfig}:
```

```
makeJob {
  name = "sshd";
  job = "
    description \"SSH server\"
    start on network-interfaces/started

    start script
      if ! test -f /etc/ssh/ssh_host_dsa_key; then
        ${openssh}/bin/ssh-keygen ...
      fi
      ...
    end script

    respawn ${openssh}/sbin/sshd -D -f ${sshdConfig}
  ";
}
```


The top-level system configuration

`/etc/nixos/nixos/system/system.nix`

- ▶ Top-level Nix expression.
- ▶ Calls other expressions to build the Upstart jobs, kernel, initrd, boot scripts...
- ▶ Takes as argument a *system configuration* — nested attribute set specifying system parameters.

The system configuration file

/etc/nixos/configuration.nix

```
{
  boot = { grubDevice = "/dev/hda"; };
  fileSystems = [
    { mountPoint = "/";
      device = "/dev/hda1";
    }
  ];
  swapDevices = [ { device = "/dev/hdb1"; } ];
  services = {
    sshd = {
      enable = true;
      forwardX11 = true;
    };
  };
}
```

Building it

Normal operation: switch

```
$ nixos-rebuild switch
```

Builds configuration (by calling `system.nix` with `/etc/nixos/configuration.nix` as argument), makes it boot default, activates it. Previous configurations still reachable through boot menu.

Test

```
$ nixos-rebuild test
```

Builds and activates configuration. Reboot reverts.

Building it

Normal operation: switch

```
$ nixos-rebuild switch
```

Builds configuration (by calling `system.nix` with `/etc/nixos/configuration.nix` as argument), makes it boot default, activates it. Previous configurations still reachable through boot menu.

Test

```
$ nixos-rebuild test
```

Builds and activates configuration. Reboot reverts.

Booting it

Grub boot menu is synthesized from the available (non-garbage collected) generations of the system profile.

```
GNU GRUB  version 0.97  (639K lower / 130048K upper memory)
```

```
NixOS - Default
```

```
NixOS - Configuration 27 (2007-04-16 16:18:13)
```

```
NixOS - Configuration 26 (2007-04-16 16:17:40)
```

```
NixOS - Configuration 25 (2007-04-16 15:31:26)
```

```
NixOS - Configuration 24 (2007-04-04 19:09:46)
```

```
NixOS - Configuration 23 (2007-04-04 19:06:14)
```

```
NixOS - Configuration 22 (2007-04-04 11:53:53)
```

```
NixOS - Configuration 21 (2007-03-01 16:36:00)
```

```
NixOS - Configuration 20 (2007-02-28 14:20:35)
```

```
NixOS - Configuration 19 (2007-02-28 10:35:47)
```

```
NixOS - Configuration 18 (2007-02-28 02:30:04)
```

```
NixOS - Configuration 17 (2007-02-22 22:45:12)
```



Use the ↑ and ↓ keys to select which entry is highlighted.
Press enter to boot the selected OS, 'e' to edit the
commands before booting, or 'c' for a command-line.

How pure are we?

- ▶ No `/bin` (with 1 exception), `/sbin`, `/lib`, `/usr`.
 - ▶ Sole exception: `/bin/sh`
- ▶ Almost all of `/etc` resides in the Nix store
 - ▶ E.g. `sshd_config`, Upstart job specifies full store path
 - ▶ But some configuration files are cross-cutting (`/etc/resolv.conf`, `/etc/services`), so we symlink them in `/etc`
- ▶ Mutable state (`/var`): don't do anything special with it
 - ▶ Nasty: hybrid configuration / state: `/etc/passwd`

How pure are we?

- ▶ No `/bin` (with 1 exception), `/sbin`, `/lib`, `/usr`.
 - ▶ Sole exception: `/bin/sh`
- ▶ Almost all of `/etc` resides in the Nix store
 - ▶ E.g. `sshd_config`, Upstart job specifies full store path
 - ▶ But some configuration files are cross-cutting (`/etc/resolv.conf`, `/etc/services`), so we symlink them in `/etc`
- ▶ Mutable state (`/var`): don't do anything special with it
 - ▶ Nasty: hybrid configuration / state: `/etc/passwd`

How pure are we?

- ▶ No `/bin` (with 1 exception), `/sbin`, `/lib`, `/usr`.
 - ▶ Sole exception: `/bin/sh`
- ▶ Almost all of `/etc` resides in the Nix store
 - ▶ E.g. `sshd_config`, Upstart job specifies full store path
 - ▶ But some configuration files are cross-cutting (`/etc/resolv.conf`, `/etc/services`), so we symlink them in `/etc`
- ▶ Mutable state (`/var`): don't do anything special with it
 - ▶ Nasty: hybrid configuration / state: `/etc/passwd`

The downside

- ▶ X11/KDE configuration takes up 656 MiB in 236 store paths.
- ▶ That's okay, but...
- ▶ ...if Glibc changes, then we'll need another 656 MiB.
- ▶ But disk space is cheap.
- ▶ However, build / download time isn't.
- ▶ Analogous to purely functional data structures:
 - ▶ `x` : tail list is cheap
 - ▶ `init list ++ [x]` is expensive
- ▶ Binary patching helps a lot.

Conclusion

- ▶ NixOS shows that a purely functional system configuration model is feasible and practical.
- ▶ Advantages: reproducibility, determinism, predictable upgrading.
- ▶ Worth mentioning: multi-user package management, any user can install software, with sharing.
- ▶ Disadvantage: can take up to 2x as much disk space.

More information / download

- ▶ <http://www.nixos.org/>
- ▶ ISO images for x86, x86_64.
- ▶ Easy way to play with it: install in a VM or on a USB stick.