NixOS: A Purely Functional Linux Distribution ICFP 2008, Victoria, Canada

Eelco Dolstra¹ Andres Löh²

¹Delft University of Technology, EWI, Department of Software Technology

 2 Utrecht University, Department of Information and Computing Sciences

September 24, 2008





Universiteit Utrecht

Overview

- ▶ Operating systems are installed and managed using tools that have an *imperative model*
- ► This causes lots of problems: upgrading is unreliable, rollbacks are hard, etc.
- ► This paper shows that it is possible to implement a system with a *purely functional model*
- ▶ Implemented in a Linux distribution called NixOS

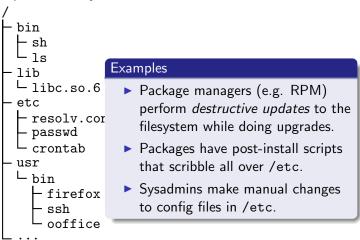
Introduction

Your operating system and applications are managed in an **imperative way**.

```
└ libc.so.6
etc
  resolv.conf
passwd
  crontab
usr
└ bin
    firefox
    ssh
    ooffice
```

Introduction

Your operating system and applications are managed in an **imperative way**.



Why is statefulness bad?

- No predictability (determinism)
 - ▶ If an action depends on an ill-defined initial state, then the result is probably ill-defined
 - ▶ This is why upgrading is riskier than a full re-install
- ► Configuration actions clobber the previous configuration
 - No multiple versions
 - No atomicity
 - No rollbacks
 - Hard to safely test a configuration
- No traceability
 - Configuration is the result of a sequence of (sometimes manual) imperative actions over time
 - Hard to reproduce a configuration

Why is statefulness bad?

- No predictability (determinism)
 - ▶ If an action depends on an ill-defined initial state, then the result is probably ill-defined
 - ▶ This is why upgrading is riskier than a full re-install
- Configuration actions clobber the previous configuration
 - No multiple versions
 - No atomicity
 - No rollbacks
 - Hard to safely test a configuration
- ► No traceability
 - Configuration is the result of a sequence of (sometimes manual) imperative actions over time
 - Hard to reproduce a configuration

Why is statefulness bad?

- No predictability (determinism)
 - ► If an action depends on an ill-defined initial state, then the result is probably ill-defined
 - ▶ This is why upgrading is riskier than a full re-install
- ► Configuration actions clobber the previous configuration
 - No multiple versions
 - No atomicity
 - No rollbacks
 - Hard to safely test a configuration
- No traceability
 - Configuration is the result of a sequence of (sometimes manual) imperative actions over time
 - Hard to reproduce a configuration

There is a better way!

- ► Solution: **borrow from purely functional languages**.
- ▶ We're going to...
 - Build packages from pure functions
 - Store packages as immutable values in the filesystem
 - Extend this to the other static bits of the system

There is a better way!

- ► Solution: borrow from purely functional languages.
- ▶ We're going to...
 - ▶ Build packages from pure functions
 - Store packages as immutable values in the filesystem
 - Extend this to the other static bits of the system

Old stuff: Nix

There is a better way!

- ► Solution: borrow from purely functional languages.
- ▶ We're going to...
 - ▶ Build packages from pure functions
 - Store packages as immutable values in the filesystem
 - Extend this to the other static bits of the system

Old stuff: Nix

New stuff: NixOS

Nix: Purely functional package management

Nix is a purely functional package manager.

- Lazy, higher-order, dynamically typed, purely functional language to describe how to build and compose packages.
- Build results only depend on declared inputs.
- Packages are stored as a kind of purely functional data structure:
 - ► They never change after they have been built
 - ▶ So in particular they're never overwritten by newer versions
 - Unused packages are garbage collected automatically.

Nix: Purely functional package management

Nix is a purely functional package manager.

- ► Lazy, higher-order, dynamically typed, purely functional language to describe how to build and compose packages.
- ▶ Build results only depend on declared inputs.
- ► Packages are stored as a kind of purely functional data structure:
 - ▶ They never change after they have been built.
 - ▶ So in particular they're never overwritten by newer versions.
 - Unused packages are garbage collected automatically.

Nix: Purely functional package management

Nix is a purely functional package manager.

- Lazy, higher-order, dynamically typed, purely functional language to describe how to build and compose packages.
- ▶ Build results only depend on declared inputs.
- Packages are stored as a kind of purely functional data structure:
 - ▶ They never change after they have been built.
 - ▶ So in particular they're never overwritten by newer versions.
 - Unused packages are garbage collected automatically.

Nix store

```
Central idea: store all packages in isolation from each other:
```

```
/nix/store/axrzx0rh0ivw...
-firefox-2.0.0.3
```

Paths contain a 160-bit **cryptographic hash** of **all** inputs used to build the package:

- Sources
- Libraries
- Compilers
- Build scripts

```
/nix/store
  19w6773m1msy...-openssh-4.6p1
     bin
     ∟ ssh
     sbin
     ∟ sshd
  smkabrbibqv7...-openss1-0.9.86
  ∟<sub>lib</sub>
      └ libssl.so.0.9.8
  c6jbqm2mc0a7...-zlib-1.2.3
  ∟<sub>lib</sub>
     \sqsubseteq libz.so.1.2.3
  im276akmsrhv...-glibc-2.5
  ∟<sub>lib</sub>
      └ libc.so.6
```

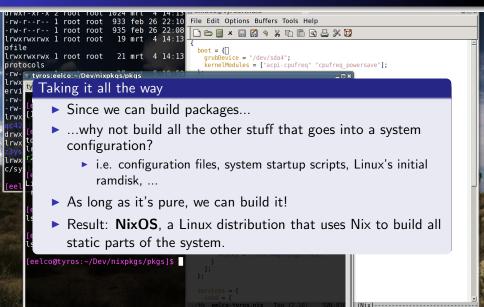
Nix expressions

```
Nix expressions describe how to build packages.
{stdenv, fetchurl, openssl, zlib}:
stdenv.mkDerivation {
  name = "openssh-4.6p1";
  src = fetchurl {
    url = http://.../openssh-4.6p1.tar.gz;
    sha256 = "Ofpjlr3bfind0y94bk442x2p...";
  };
  buildCommand = ''
    tar xjf $src
    ./configure --prefix=$out --with-openssl=${openssl}
    make; make install
  , , ;
```

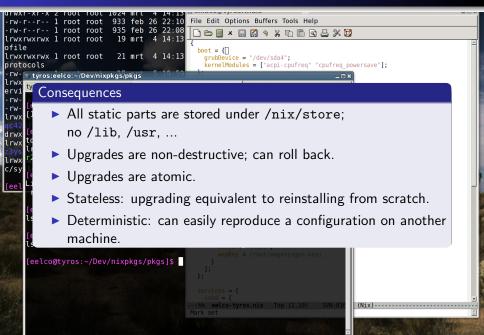
Nix expressions

```
Nix expressions describe how to build packages.
{stdenv, fetchurl, openssl, zlib}:
stdenv.mkDerivation {
  name = "openss
                  Evaluating the result of this function will pro-
  src = fetchur]
                  duce an OpenSSH package in the Nix store.
    url = http:/
                  /nix/store
    sha256 = "01
                     19w6773m1msy...-openssh-4.6p1
  buildCommand =
                        bin
    tar xjf $src
                         L<sub>ssh</sub>
    ./configure
                         sbin
    make; make i
                         L sshd
  ,,;
```

NixOS



NixOS



NixOS

```
933 feb 26 22:10 File Edit Options Buffers Tools Help
           1 root root
-rw-r--r-- 1 root root 935 feb 26 22:08
                                           lrwxrwxrwx 1 root root 19 mrt 4 14:13
lofile
                                            boot = {
lrwxrwxrwx 1 root root
                         21 mrt 4 14:13
                                             grubDevice = "/dev/sda4";
protocols
                                             kernelModules = ["acpi-cpufreg" "cpufreg powersave"]:
- TW- tyros:eelco:~/Dev/nixpkgs/pkgs
    tyros:eelco:~/Dev/nixpkqs/p... x tyros:eelco:~/Dev/nixpkqs/p... x tyros:eelco:~/Dev/nixpkqs/p... x
ervi
- rw-
    [eelco@tyros:~/Dev/nixpkgs/pkgs]$ emacs /etc/nixos/configuration.nix &
    [1] 19510 oot root 57 mrt 4 02 26 { mounts
l rwx
     [eelco@tyros:~/Dev/nixpkgs/pkgs]$ ls -l /bin/
d rwx
    total 0
1 rwx
     lrwxrwxrwx 1 root root 59 mrt 4 14:13 sh -> /nix/store/vlmni7n98vz3lrp
    r2gp8ng8iw29sbcgs-bash-3.2/bin/sh
lrwx
c/sy
     [eelco@tyros:~/Dev/nixpkgs/pkgs]$ cat /proc/version
     Linux version 2.6.20-skas3-v9-pre9-default (nix@scratchy) (collect2: ld
     returned 1 exit status) #1 SMP Thu Feb 8 17:23:43 UTC 2007
     [eelco@tyros:~/Dev/nixpkgs/pkgs]$ ls -le/usrng = {
     ls: cannot access /usr: No such file or directory
     [eelco@tyros:~/Dev/nixpkgs/pkgs]$ ls -l /lib
     ls: cannot access /lib: No such file or directory
     [eelco@tyros:~/Dev/nixpkgs/pkgs]$
```

```
Nix expression for ssh_config
{config, pkgs}:
pkgs.writeText "ssh_config" ''
  SendEnv LANG LC_ALL ...
  ${if config.services.sshd.forwardX11 then ''
    ForwardX11 yes
    XAuthLocation ${pkgs.xorg.xauth}/bin/xauth
  '' else ''
    ForwardX11 no
  ,,}
, ,
```

```
Nix expression for ssh_config
{config, pkgs}:
pkgs.writeText "ssh_config" ''
  SendEnv LANG LC_ALL ...
  ${if config.services.sshd.forwardX11 then
    ForwardX11 yes
    XAuth ocation ${pkgs.xorg.xauth}/bin/xauth
  ', else
       Laziness in action!
, ,
```

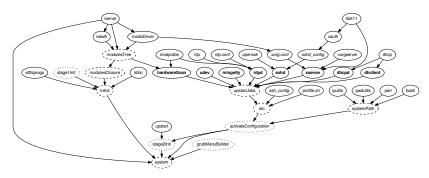
```
Nix expression for ssh_config
{config, pkgs}:
pkgs.writeText "ssh_config"
  SendEnv LANG LC_ALL ...
  ${if config.serv Nix store
    ForwardX11 yes
                     /nix/store
    XAuthLocation
                     - 331cnh62y113...-ssh_config
  '' else ''
                      - kyv6n69a40q6...-xauth-1.0.2
    ForwardX11 no
                       L bin
  ,,}
                          ∟ xauth
, ,
```

```
Nix expression for ssh_config
 {config, pkgs}:
 pkgs.writeText "ssh_config"
   SendEnv LANG LC_ALL ...
   ${if config.serv Nix store
     ForwardX11 yes /nix/store
     XAuthLocation :

→ 331cnh62yll3...-ssh_config
   '' else ''
                        kvv6n69a40q6...-xauth-1.0.2
     ForwardX11 no
Generated file: 33lcnh62yll3...-sshd_config
SendEnv LANG LC ALL ...
ForwardX11 yes
XAuthLocation /nix/store/kyv6n69a40q6...-xauth-1.0.2/bin/xauth
```

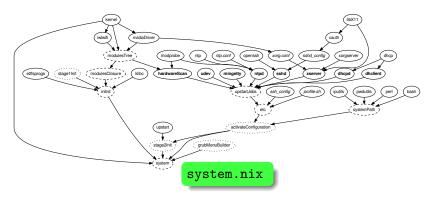
NixOS build time dependency graph

Nix expressions to build each part of the system: system packages, applications, their dependencies, kernel modules, initrd, configuration files, Upstart jobs, boot scripts, ...



NixOS build time dependency graph

Nix expressions to build each part of the system: system packages, applications, their dependencies, kernel modules, initrd, configuration files, Upstart jobs, boot scripts, ...



The system configuration file

```
/etc/nixos/configuration.nix
  boot = { grubDevice = "/dev/sda"; };
  fileSystems = [
    { mountPoint = "/";
      device = "/dev/sda1";
 ];
  swapDevices = [ { device = "/dev/sdb1"; } ];
  services = {
    sshd = {
      enable = true;
      forwardX11 = true;
```

The system configuration file

```
/etc/nixos/configuration.nix
  boot = { grubDevice = "/dev/sda"; };
  fileSystems = [
    { mountPoint = "/";
      device = "/dev/sda1";
 1:
  swapDevices = [
  services = {
    sshd = {
      enable = tru
      forwardX11 =
```

End-user perspective

- Edit configuration.nix.
- Run nixos-rebuild.
- ► This builds system.nix and runs its activation script.
- Non-destructive; various rollback mechanisms.

NixOS — Grub boot menu

```
NixOS - Configuration 128 (2008-09-17 14:36:01 - 2.6.26.5-default)
NixOS - Configuration 127 (2008-09-17 14:35:01 - 2.6.26.5-default)
NixOS - Configuration 126 (2008-09-17 09:54:34 - 2.6.26.5-default)
NixOS - Configuration 125 (2008-09-16 23:26:05 - 2.6.25.16-default)
NixOS - Configuration 124 (2008-09-16 23:26:05 - 2.6.25.16-default)
NixOS - Configuration 123 (2008-09-02 14:41:24 - 2.6.25.16-default)
NixOS - Configuration 123 (2008-09-02 11:48:24 - 2.6.26.3-default)
NixOS - Configuration 122 (2008-09-02 11:40:04 - 2.6.26.3-default)
NixOS - Configuration 121 (2008-08-29 14:08:21 - 2.6.26.3-default)
NixOS - Configuration 120 (2008-08-29 11:09:07 - 2.6.26.3-default)
NixOS - Configuration 120 (2008-08-29 11:09:07 - 2.6.26.3-default)
NixOS - Configuration 120 (2008-08-29 11:09:07 - 2.6.26.3-default)
```

Ţ

Use the \uparrow and \downarrow keys to select which entry is highlighted. Press enter to boot the selected OS, 'e' to edit the commands before booting, or 'c' for a command-line.

NixDS - Configuration 118 (2008-08-22 09:52:55 - 2.6.26.3-default) NixDS - Configuration 117 (2008-08-21 15:31:39 - 2.6.26.3-default)

GNU GRUB version 0.97 (639K lower / 129984K upper memory)

GNU/Linux

Code

- ▶ No /bin, /sbin, /lib, /usr.
- ▶ Only exception: /bin/sh
- Only minor changes needed to a few packages.

```
[eelco@dutibo:~]$ ls -l /bin/
total 1
lrwxrwxrwx 1 root root 63 2008-09-24 02:09 sh -> /nix/store/a46s354jbdm]
[eelco@dutibo:~]$ ls -l /usr
ls: cannot access /usr: No such file or directory
[eelco@dutibo:~]$ ls -l /lib
ls: cannot access /lib: No such file or directory
```

Configuration data

- Almost all of /etc resides in the Nix store
- E.g. sshd_config, Upstart job specifies full store path.
- But some configuration files are cross-cutting (/etc/nsswitch.conf, /etc/services), so we symlink them in /etc.

Mutable state

- E.g. /var/log.
- ▶ Handled normally; not part of the purely functional model.
- Nasty: hybrid configuration+state: /etc/passwd; handled by the imperative activation script.

```
[eelco@dutibo:~]$ ls -l /var/
total 11
drwxr-xr-x 6 root root 1024 2008-09-24 17:10 cache
drwx--x--- 3 root root 1024 2008-01-07 14:52 cron
drwxr-xr-x 4 root root 1024 2008-06-06 10:20 db
dr-xr-xr-x 2 root root 1024 2008-01-07 14:52 empty
drwxr-xr-x 7 root root 1024 2008-06-06 10:48 lib
drwx------ 3 root root 1024 2008-01-07 14:52 lock
drwxr-xr-x 5 root root 1024 2008-09-23 03:15 log
drwxr-xr-x 8 root root 1024 2008-09-24 02:09 run
```

Build actions

- ▶ Can't guarantee that build actions are pure.
- ► E.g. output could use
 - System time
 - ▶ The network
 - Files outside the Nix store
- Experiment: built 485 packages on two machines.
- ▶ Result: out of 165 927 files, 5059 differed.
- Almost all (possibly all) are due to timestamps or the hostname.

Conclusion

- Applied lazy, purely functional paradigm to a very different domain — system configuration management.
 - Purity and laziness are incredibly useful here.
- ▶ NixOS shows that a purely functional system configuration model is feasible and practical.
- ► Advantages: reproducibility, predictable upgrading, rollbacks, multi-user package management, ...

More information / download

- http://nixos.org/
- ► ISO images for x86, x86_64 are available.