

NixOS: The Purely Functional Linux Distribution

Eelco Dolstra
LogicBlox

21 November 2013



So what was wrong with the previous 298 distributions?

So what was wrong with the previous 298 distributions?

- Upgrading is risky business
 - ▶ No way to roll back
 - ▶ Not atomic / transactional

So what was wrong with the previous 298 distributions?

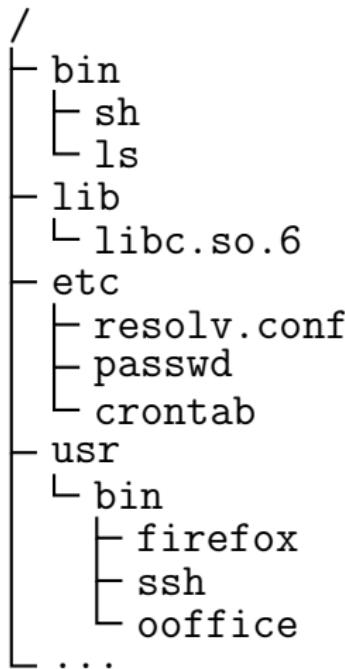
- Upgrading is risky business
 - ▶ No way to roll back
 - ▶ Not atomic / transactional
- Not declarative
 - ▶ Configuration is result of automated/manual changes over time
 - ▶ Therefore hard to reproduce

So what was wrong with the previous 298 distributions?

- Upgrading is risky business
 - ▶ No way to roll back
 - ▶ Not atomic / transactional
- Not declarative
 - ▶ Configuration is result of automated/manual changes over time
 - ▶ Therefore hard to reproduce
- Package management limitations
 - ▶ Hard to have multiple versions of a package
 - ▶ Need to be root to install packages
 - ▶ Upgrading one package may break others (“DLL hell”)

Analysis

Conventional Linux distributions have an **imperative** (stateful) approach to configuration / package management.



Analysis

Conventional Linux distributions have an **imperative** (stateful) approach to configuration / package management.

```
/  
└ bin  
    └ sh  
    └ ls  
- lib  
    └ libc.so.6  
- etc  
    └ resolv.conf  
    └ passwd  
    └ crontab  
- usr  
    └ bin  
        └ firefox  
        └ ssh  
        └ ooffice  
...  
...
```

Examples

- Package managers like RPM perform *destructive updates* to the filesystem while doing upgrades.
- Packages have post-install scripts that scribble all over /etc.
- Sysadmins make manual changes to config files in /etc.

Enter NixOS

NixOS has a purely functional model

(Based on Nix, the purely functional package manager.)

This means:

- Packages, config files and other static parts of the system are built by *pure functions*.
- They are *immutable* after they have been built.

Consequences

- All static parts are stored under `/nix/store`;
(almost) no `/bin`, `/lib`, `/usr`, ...
- Upgrades are non-destructive; can roll back.
- Upgrades are atomic.
- Stateless: upgrading equivalent to reinstalling from scratch.
- Deterministic: can easily reproduce a configuration on another machine.

How it works: the Nix store

Nix stores all packages in isolation from each other:

/nix/store/rpdqxnnilb0cg...
-firefox-25.0

Paths contain a 160-bit **cryptographic hash** of all inputs used to build the package:

- Sources
- Libraries
- Compilers
- Build scripts
- ...

```
/nix/store
└─ 19w6773m1msy...-openssh-4.6
    └─ bin
        └─ ssh
            └─ sbin
                └─ sshd
└─ smkabrbibqv7...-openssl-0.9.
    └─ lib
        └─ libssl.so.0.9.8
└─ c6jbqm2mc0a7...-zlib-1.2.3
    └─ lib
        └─ libz.so.1.2.3
└─ im276akmsrhv...-glibc-2.5
    └─ lib
        └─ libc.so.6
```

NixOS is declarative

```
/etc/nixos/configuration.nix
```

```
{  
    boot.loader.grub.bootDevice = "/dev/sda";  
    fileSystems."/".device = "/dev/sda1";  
    swapDevices = [ { device = "/dev/sdb1"; } ];  
    services.sshd.enable = true;  
    services.sshd.forwardX11 = true;  
}
```

NixOS is declarative

To enable Apache, add this to configuration.nix:

```
services.httpd.enable = true;  
services.httpd.documentRoot = "/bla";
```

NixOS is declarative

To enable Apache, add this to configuration.nix:

```
services.httpd.enable = true;  
services.httpd.documentRoot = "/bla";
```

Or to enable X11/KDE:

```
services.xserver.enable = true;  
services.xserver.displayManager.kde4.enable = true;
```

NixOS is declarative

To enable Apache, add this to configuration.nix:

```
services.httpd.enable = true;  
services.httpd.documentRoot = "/bla";
```

Or to enable X11/KDE:

```
services.xserver.enable = true;  
services.xserver.displayManager.kde4.enable = true;
```

Or to make Firefox available to users:

```
environment.systemPackages = [ pkgs.firefox ];
```

NixOS is declarative

To enable Apache, add this to configuration.nix:

```
services.httpd.enable = true;  
services.httpd.documentRoot = "/bla";
```

Or to enable X11/KDE:

```
services.xserver.enable = true;  
services.xserver.displayManager.kde4.enable = true;
```

Or to make Firefox available to users:

```
environment.systemPackages = [ pkgs.firefox ];
```

Then run:

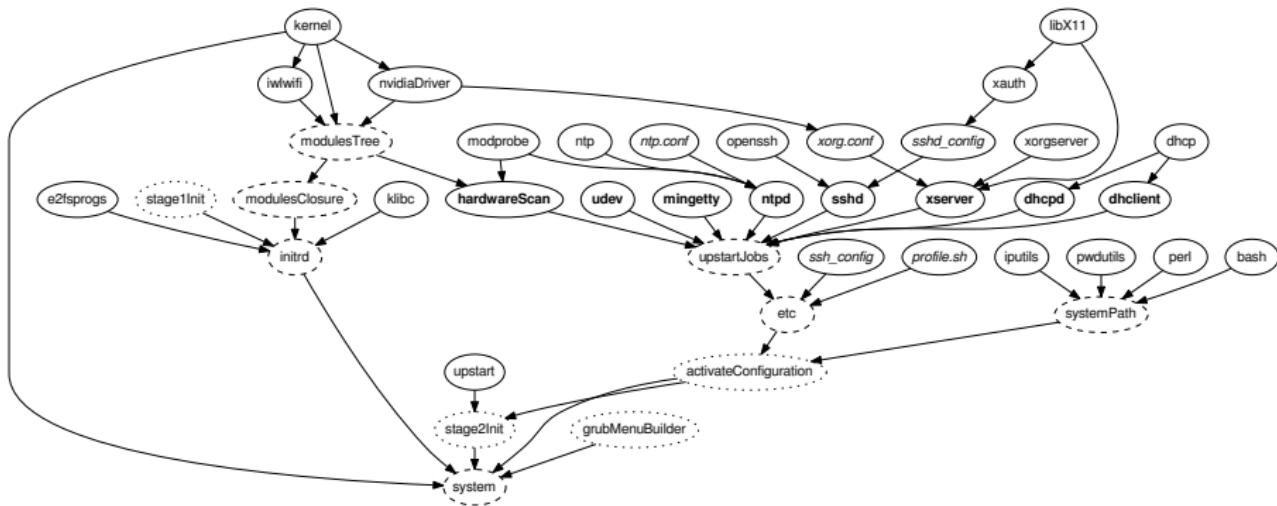
```
$ nixos-rebuild switch
```

Reproducibility

`nixos-rebuild switch` invokes Nix to build or download the entire system configuration: packages, boot scripts, config files, ...

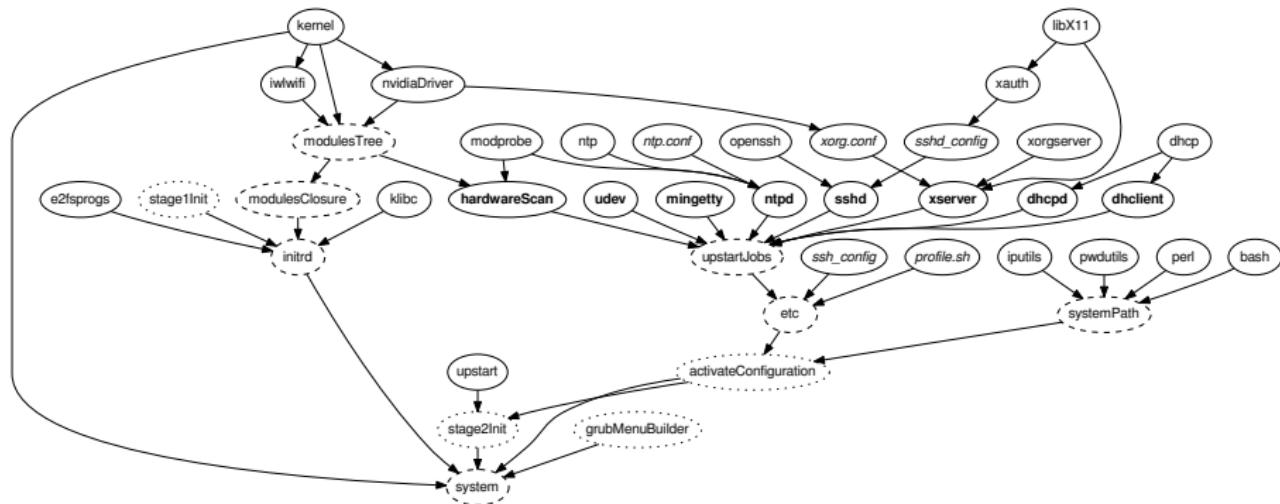
Reproducibility

`nixos-rebuild switch` invokes Nix to build or download the entire system configuration: packages, boot scripts, config files, ...



Reproducibility

`nixos-rebuild switch` invokes Nix to build or download the entire system configuration: packages, boot scripts, config files, ...



Since this is purely functional, configurations can always be reproduced.

Rollbacks are easy

To undo the last upgrade:

```
$ nixos-rebuild switch --rollback
```

Rollbacks are easy

To undo the last upgrade:

```
$ nixos-rebuild switch --rollback
```

Alternatively, revert configuration.nix, e.g. remove

```
services.httpd.enable = true;  
services.httpd.documentRoot = "/bla";
```

and run

```
$ nixos-rebuild switch
```

Rollbacks are really easy

GNU GRUB version 2.00

- NixOS - Configuration 312 (2013-11-20 - 13.10.35455.45219b9)
- NixOS - Configuration 311 (2013-11-20 - 13.10.35455.45219b9)
- NixOS - Configuration 310 (2013-11-19 - 13.10.35450.912f584)
- NixOS - Configuration 309 (2013-11-19 - 13.10.35450.912f584)
- NixOS - Configuration 308 (2013-11-19 - 13.10.35450.912f584)
- NixOS - Configuration 307 (2013-11-08 - 13.10.35442.b77a2cd)
- NixOS - Configuration 306 (2013-11-02 - 13.10.35432.a6cc1e4)
- NixOS - Configuration 305 (2013-11-01 - 13.10.35427.6fd96b)
- NixOS - Configuration 304 (2013-10-31 - 13.10pre35379.70a2c54)
- NixOS - Configuration 303 (2013-10-30 - 13.10pre35354.8150f1a)
- NixOS - Configuration 302 (2013-10-25 - 13.10pre35223.ab94ccf)
- NixOS - Configuration 301 (2013-10-25 - 13.10pre35209.6f911ed)
- NixOS - Configuration 300 (2013-10-23 - 13.10pre35155.897329f)
- NixOS - Configuration 299 (2013-10-23 - 13.10pre35155.897329f)
- NixOS - Configuration 298 (2013-10-21 - 13.10pre35085.81ef604)
- NixOS - Configuration 297 (2013-10-16 - 13.10pre34974.c8f261c)

Use the ↑ and ↓ keys to select which entry is highlighted.
Press enter to boot the selected OS, 'e' to edit the commands
before booting or 'c' for a command-line. ESC to return

Testing a configuration is painless

To test a new configuration without making it the boot default:

```
$ nixos-rebuild test
```

Testing a configuration is painless

To test a new configuration without making it the boot default:

```
$ nixos-rebuild test
```

To build a VM containing the new configuration:

```
$ nixos-rebuild build-vm
```

```
[eelco@hagbard:~/Dev/nixos]$ nixos-rebuild build-vm --fast
building the system configuration...
trace: Obsolete option `services.sshd.enable' is defined instead of `services.openssh.enable'.
Done. The virtual machine can be started by running ./result/bin/run-hagbard-vm.
```

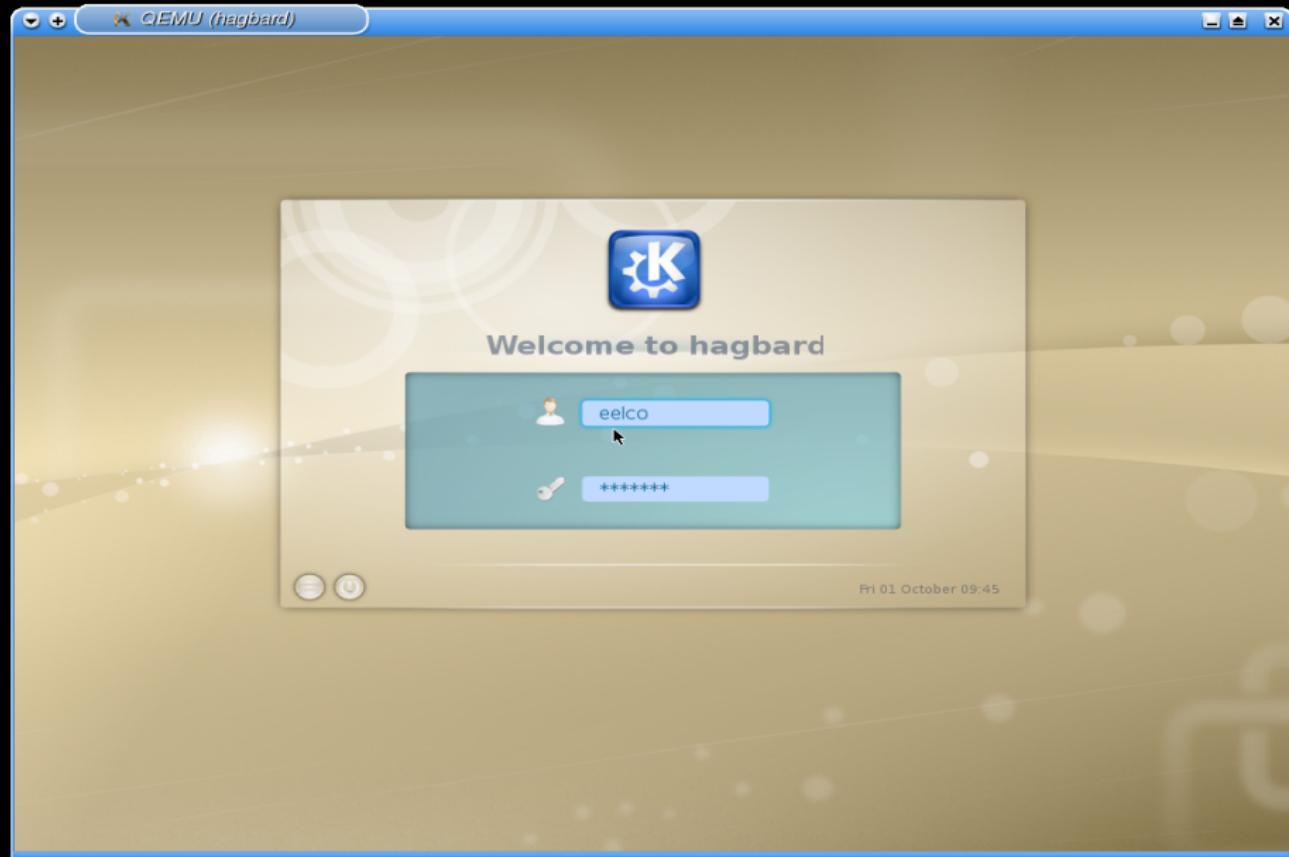
```
[eelco@hagbard:~/Dev/nixos]$ ./result/bin/run-hagbard-vm
```

QEMU (hagbard)

```
[ 0.659344] virtio-pci 0000:00:03.0: PCI INT A -> Link[LNK0] -> GSI 11 (level, high) -> IRQ 11
[ 0.671629] ACPI: PCI Interrupt Link [LNKD] enabled at IRQ 10
[ 0.672124] virtio-pci 0000:00:04.0: PCI INT A -> Link[LNKD] -> GSI 10 (level, high) -> IRQ 10
[ 0.679294] vda: unknown partition table
[ 0.682585] scsi0 : ata_piix
[ 0.684051] scsi1 : ata_piix
[ 0.698057] ata1: PATA max MWDMA2 cmd 0x1f0 ctl 0x3f6 bmdma 0xc000 irq 14
[ 0.699059] ata2: PATA max MWDMA2 cmd 0x170 ctl 0x376 bmdma 0xc008 irq 15
[ 0.852299] ata2.00: ATAPI: QEMU DVD-ROM, 0.12.5, max UDMA/100
[ 0.855273] ata2.00: configured for MWDMA2
[ 0.857502] scsi 1:0:0:0: CD-ROM           QEMU      QEMU DVD-ROM    0.12 PQ: 0 ANSI: 5
[ 0.870604] sr0: scsi3-mmc drive: 4x/4x xa/form2 tray
[ 0.871193] Uniform CD-ROM driver Revision: 3.20
starting device mapper and LVM...
Reading all physical volumes. This may take a while...
No volume groups found
No volume groups found
[ 0.944976] Switching to clocksource hpet
failed to resume...
mounting /dev/vda on /...
fsck from util-linux-ng 2.17.2
/nix/store/6vn4la0b98gz18wynsbamz/jxwby4yjgn-extra-utils/bin/fsck.ext3 (1) -- /dev/vda1 fsck.ext3
-a -C0 /dev/vda
/dev/vda: recovering journal
/dev/vda: clean, 1322/32768 files, 21486/131072 blocks
[ 1.007550] kjournald starting. Commit interval 5 seconds
[ 1.019665] EXT3 FS on vda, internal journal
[ 1.020993] EXT3-fs: mounted filesystem with writeback data mode.
mounting //10.0.2.4/qemu on /hostfs...
[ 1.032796] Slow work thread pool: Starting up
[ 1.033430] Slow work thread pool: Ready
mounting /mnt-root/hostfs/nix/store on /nix/store...
<<< NixOS Stage 2 >>>
running activation script...
setting up /etc...
```

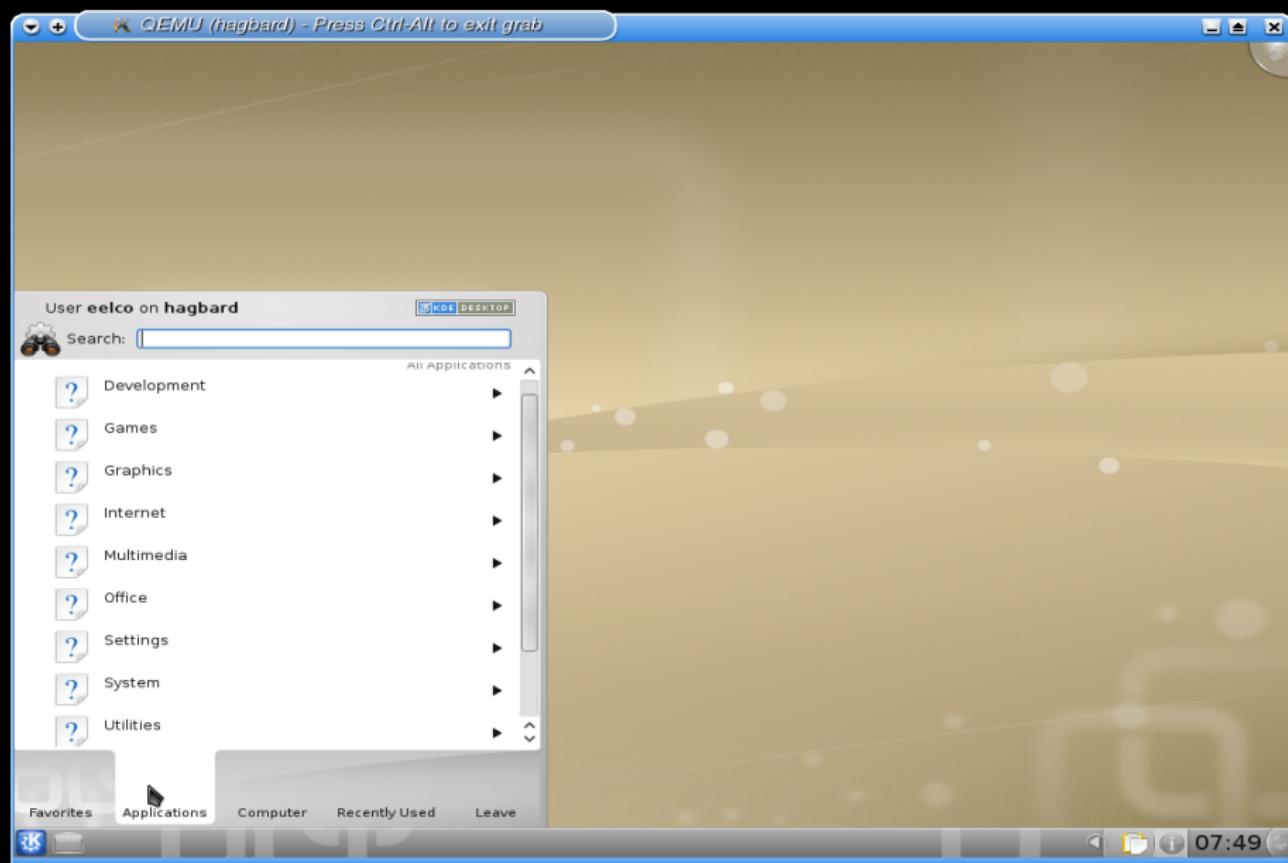
```
[eelco@hagbard:~/Dev/nixos]$ nixos-rebuild build-vm --fast
building the system configuration...
trace: Obsolete option `services.sshd.enable' is defined instead of `services.openssh.enable'.
Done. The virtual machine can be started by running ./result/bin/run-hagbard-vm.
```

```
[eelco@hagbard:~/Dev/nixos]$ ./result/bin/run-hagbard-vm
```



```
[eelco@hagbard:~/Dev/nixos]$ nixos-rebuild build-vm --fast
building the system configuration...
trace: Obsolete option `services.sshd.enable' is defined instead of `services.openssh.enable'.
Done. The virtual machine can be started by running ./result/bin/run-hagbard-vm.
```

```
[eelco@hagbard:~/Dev/nixos]$ ./result/bin/run-hagbard-vm
```



Multi-user package management

Non-root users can install software:

```
alice$ nix-env -i firefox-25.0.1
```

```
alice$ firefox --version
```

```
Mozilla Firefox 25.0.1
```

Multi-user package management

Non-root users can install software:

```
alice$ nix-env -i firefox-25.0.1
```

```
alice$ firefox --version
```

```
Mozilla Firefox 25.0.1
```

Without interfering with each other:

```
bob$ firefox
```

```
firefox: No such file or directory
```

```
bob$ nix-env -i firefox-3.6.27
```

```
bob$ firefox --version
```

```
Mozilla Firefox 3.6.27
```

NixOps: Declarative Provisioning and Deployment

NixOps extends the NixOS approach to networks of machines

```
{  
    database =  
        { deployment.targetEnv = "virtualbox";  
            services.postgresql.enable = true;  
        };  
  
    webserver =  
        { deployment.targetEnv = "virtualbox";  
            services.httpd.enable = true;  
            services.httpd.documentRoot = ...;  
        };  
}
```

charone: bash

```
[eelco@darkard:~/Dev/charon/examples]$ charon -s ./apache-vbox.json deploy
creating VirtualBox VM 'backend1'...creating VirtualBox VM 'backend2'...
creating VirtualBox VM 'proxy'...
Virtual machine 'charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend1' is created and registered.
UUID: 5f85392e-8f42-4b02-96a3-e28c790f91ab
Virtual machine 'charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend2' is created and registered.
UUID: c54af12c-c2c1-4fbc-84d3-943cfdd485
Settings file: '/home/eelco/VirtualBox VMs/charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend2.vbox'
Virtual machine 'charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-proxy' is created and registered.
UUID: 96d1b5d-78c5-40cd-85ca-b835a6c2d6c6
Settings file: '/home/eelco/VirtualBox VMs/charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-proxy/charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-proxy.vbox'
8%...8%...10%...20%...10%...20%...30%...40%...30%...40%...40%...50%...50%...50%...60%...60%...60%...70%...70%...80%...80%...90%...90%...80%...90%...100%
100%
100%
Clone hard disk created in format 'VDI'. UUID: ffc13a76-c415-47ae-a18f-e6e63ef45c77
Clone hard disk created in format 'VDI'. UUID: 8dc0e019-eaac-4729-a211-e99ca9a43a12
Clone hard disk created in format 'VDI'. UUID: 3bcdb8624-c13-4eee-b64d-00e28430c0b9
Waiting for VM "charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-proxy" to power on...
Waiting for VM "charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend1" to power on...
Waiting for VM "charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend2" to power on...
VM "charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend2" has been successfully started.
VM "charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend1 [Running] - Oracle VM Virtu...
waiting for IP
Machine View Devices Help
<div>
</div>
</div>
<div id="footer">
  <ul class="short-menu" id="bottom-menu">
    <li><a href="#">About us</a></li>
  </ul>
  <div id="last-modified">
    <p>This page was last updated on 2011-03-31 at 10:40:30Z by <tt>eelco</tt> (revision 26622).</p>
  </div>
</div>
</body>
</html>

[root@backend1:~]# status httpd
httpd start/running, process 2892

[root@backend1:~]# _
```

PC Speaker as /devices/platform/pcspkr/input/input2
 s rtc_cmos: rtc core: registered rtc_cmos as rtc0
 alarm up to one day, 114 bytes nvram
 lmPS-2 Generic Explorer Mouse as /devices/platform/i0042
 irtalbox main process (1226) terminated with status 1
 irtalbox main process ended, respawning
 Power Button as /devices/LNXSYSTM:00/LNKPWRDN:00/input/input0
 Power Button (PWRF)
 Sleep Button as /devices/LNXSYSTM:00/LNKSLPDN:00/input/input1
 .lpre3926-33924 (x86_64) - Linux 3.2.16 (tty1) >>>

.lpre3959-33963 (x86_64) - Linux 3.2.16 (tty1) >>>

filesystems...

charon-b39f38c6-6001-11e1-b3ff-c3e818a3331-backend2 [Running] - Oracle VM Virtu...
Machine View Devices Help
low) -> IRQ 9
5.4478011 rtc_cmos rtc...
to one day, 114 bytes nvram
an device as /devices/virtual/input/input1
r as /devices/platform/pcspkr/input/input2
00-07.0: SMBus base address uninitialized - upgr...
dr
0, IRQ 9, I/O port d820, MMIO at 0000000000000000
Generic Explorer Mouse as /devices/platform/i0042

mixos login:

mixos login:

backend2 login: mounting filesystems...

Deploying to EC2

```
{  
    database =  
        { deployment.targetEnv = "ec2";  
            deployment.ec2.region = "us-east-1";  
            services.postgresql.enable = true;  
        };  
  
    webserver =  
        { deployment.targetEnv = "ec2";  
            deployment.ec2.region = "eu-west-1";  
            services.httpd.enable = true;  
            services.httpd.documentRoot = ...;  
        };  
}
```

NixOps — Deploying

```
$ nixops create -n foo ./network.nix  
$ nixops deploy -d foo
```

This will:

- Create all machines
- Build/download all dependencies
- Upload them to the machines
- Activate any necessary services

To redeploy: just run `nixops deploy` again.

Conclusion

Why NixOS is awesome

- Completely declarative
- Atomic upgrades and rollbacks
- Reproducibility
- Multi-user package management
- NixOps extends this to networks and adds provisioning

More information

- <http://nixos.org/>