

Service Configuration Management

SCM-12

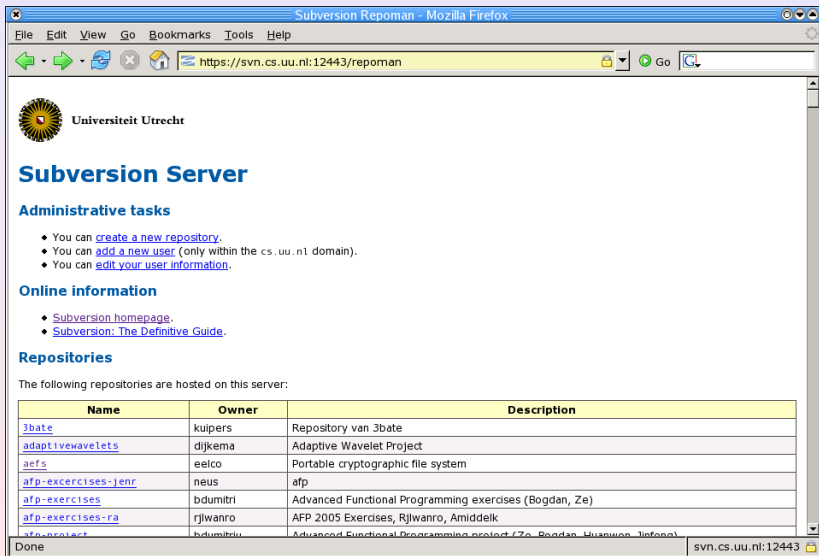
Eelco Dolstra Martin Bravenboer Eelco Visser
{eelco, martin, visser}@cs.uu.nl

Universiteit Utrecht, Faculty of Science,
Department of Information and Computing Sciences

September 6, 2005

Services: sets of running programs that provide some useful facility on a system or network.


Example: Subversion service



Subversion Repoman - Mozilla Firefox

File Edit View Go Bookmarks Tools Help

https://svn.cs.uu.nl:12443/repoman

 Universiteit Utrecht

Subversion Server

Administrative tasks

- You can [create a new repository](#).
- You can [add a new user](#) (only within the cs.uu.nl domain).
- You can [edit your user information](#).

Online information

- [Subversion homepage](#).
- [Subversion: The Definitive Guide](#).

Repositories

The following repositories are hosted on this server:

Name	Owner	Description
3bate	kuipers	Repository van 3bate
adaptivewavelets	dijkema	Adaptive Wavelet Project
aefts	eelco	Portable cryptographic file system
afp-exercises-jenr	neus	afp
afp-exercises	bdumitri	Advanced Functional Programming exercises (Bogdan, Ze)
afp-exercises-ra	rjlwanro	AFP 2005 Exercises, Rjlwanro, Amiddelk
afp-project	bdumitriu	Advanced Functional Programming project (Ze, Bogdan, Huanwen, Jinfo)

Done

svn.cs.uu.nl:12443

Example: Issue tracking service

[#NIXOS-16] NixOS should not wipe my hard drive - Stratego/XT JIRA - Mozilla Firefox

File Edit View Go Bookmarks Tools Help

https://bugs.cs.uu.nl/browse/NIXOS-16

JIRA Stratego/XT JIRA History | Log In

HOME BROWSE PROJECT FIND ISSUES CREATE NEW ISSUE QUICK SEARCH:

Issue Details [XML]

Key: NIXOS-16
Type: Bug
Status: Open
Priority: Blocker
Assignee: Armijn Hemel
Reporter: Felco Dolstra
Votes: 0
Watchers: 0

Operations

☐ Clone this issue
☐ Comment on this issue

If you were [logged in](#) you would be able to see more operations.

NixOS
NixOS should not wipe my hard drive
Created: 2005-08-23 14:26 Updated: 2005-08-23 14:26

[Return to search](#)
Issue 14 of 25 issue(s)
[<< Previous](#) | NIXOS-16 | [Next >>](#)

Component/s: None
Affects Version/s: None
Fix Version/s: None

Original Estimate:	Unknown	Remaining Estimate:	Unknown	Time Spent:	Unknown
---------------------------	---------	----------------------------	---------	--------------------	---------

Description

The installer should

a) ask for confirmation before installing
b) not wipe the existing Nix store but reuse it

[All](#) [Comments](#) [Work Log](#) [Change History](#) Sort Order:

There are no comments yet on this issue.

This site is running on Atlassian [JIRA](#) with a free **Open Source Project / Non-profit License** ([license details](#)).
[JIRA](#) is an issue and bug tracking application. [Evaluate JIRA for your organisation](#).

Done bugs.cs.uu.nl

Service deployment is hard

Service deployment involves a number of steps:

- ▶ Deploy software components (e.g., Apache, PostgreSQL, Subversion)
- ▶ Edit configuration files (e.g., **httpd.conf**, **viewcvs.conf**)
- ▶ Initialise state (e.g., logging directories, database tables)
- ▶ Start/stop processes
- ▶ ... and all of this possibly on multiple machines / platforms

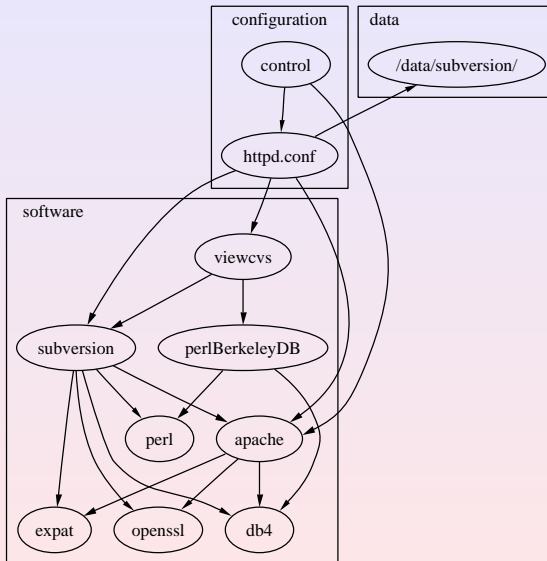
Problems

- ▶ Poor reproducibility (bad CM)
- ▶ Hard to support parallel configurations
- ▶ Cross-cutting configuration choices

Problem 1: Poor reproducibility

- ▶ Goal: it should be possible to realise a service by running a single command.
 - ▶ E.g., to move it to another machine
 - ▶ So no manual installing of missing software components, tweaking of configuration files, creating missing directories, etc.
- ▶ Why is reproducibility hard?
 - ▶ Admins often manually edit configuration files and initialise state
 - ▶ Service configuration doesn't express software component dependencies

Example



Gap between package management and service configuration

- ▶ Software components are typically deployed through package managers such as RPM
- ▶ Service configuration is typically kept under version management
- ▶ However, there is no good way to express the dependencies of the service on the software components

Problem 2: Parallel configurations

- ▶ It should be easy to create different instances of a service
 - ▶ Test vs. production servers (running on different ports, using different databases, etc.)
 - ▶ Instantiations for different users
 - ▶ Evolution through time (rollbacks)
- ▶ This is hard to support because there are typically lots of configuration files and control scripts that refer to lots of paths for components, state, static data files, etc.
 - ▶ `/etc/apache/httpd.conf`, `/etc/init.d/apache`, `/etc/apache/viewcvs.conf`, ...

Example

`/etc/apache/httpd.conf` for Subversion service (fragment)

```
ServerRoot "/var/httpd"
ServerName svn.cs.uu.nl:8080
LoadModule dav_svn_module /usr/lib/modules/mod_dav_svn.so
<Location /repos>
    AuthType Basic
    AuthDBMUserFile /data/subversion/db/svn-users
    ...
    SVNParentPath /data/subversion/repos
</Location>
ScriptAlias /viewcvs /usr/viewcvs/www/cgi/viewcvs.cgi
```

Use cases

- ▶ Try out with a different set of repositories.
- ▶ Try out a different Apache.
- ▶ Try out a different Subversion module.

Example

/etc/init.d/httpd for Subversion service (fragment)

```
/usr/sbin/apachectl -k start -f /etc/apache/httpd.conf
```

Problem 3: Cross-cutting configuration choices

- ▶ Many configuration choices are *cross-cutting*, i.e., impact many different (parts of) configuration files, scripts, etc.
- ▶ Examples:
 - ▶ Port numbers
 - ▶ Host names
 - ▶ Paths (major source of problems!)
- ▶ So a change to the configuration choices must be realised in many different places
- ▶ Lots of work
- ▶ Danger of inconsistency

Example: port number

In `/etc/init.d/httpd.conf`

```
ServerName www.example.org:12443  
Listen 12443  
<VirtualHost _default_:12443>
```

In `repoman.pl`

```
my $url = "https://www.example.org:12443/"  
print "... <a href='$url/repos/$repoName'> ...";
```

The solution

- ▶ The solution: integrate build management, software deployment, and service deployment into a single formalism
- ▶ Namely, the *Nix* deployment system
- ▶ Nix was originally created for *software deployment*
- ▶ Nice properties:
 - ▶ Automatic building of components and their dependencies
 - ▶ Side-by-side deployment, rollbacks
 - ▶ Prevention of undeclared dependencies, automatic determination of runtime dependencies
 - ▶ Functional component description language
- ▶ All these are also useful for *service deployment*
- ▶ Central idea of this paper: *treat services as components*

The Nix Deployment System

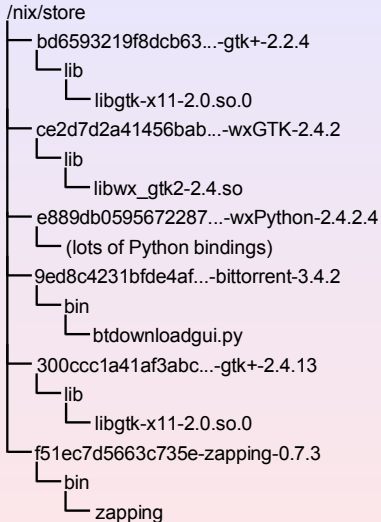
- ▶ Central idea: store all components in isolation.
- ▶ Unique paths:

```
/nix/store/605332199533e73b...-gtk+-2.2.4
```

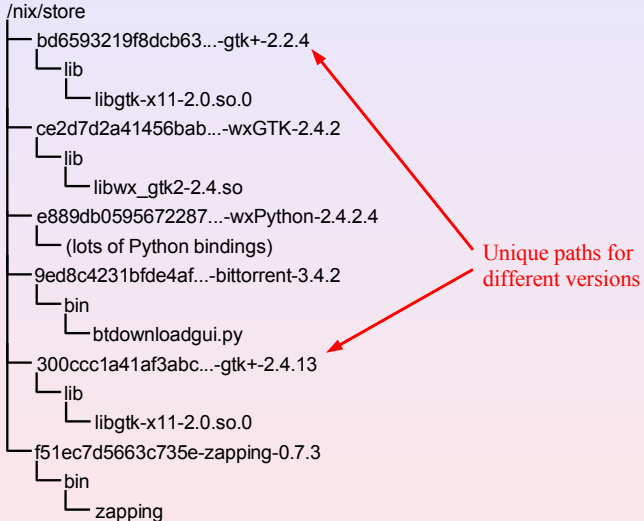
which is an SHA-256 hash of **all** inputs used to build the component:

- ▶ Sources
 - ▶ Libraries
 - ▶ Compilers
 - ▶ Build scripts
 - ▶ Build parameters
 - ▶ System type
 - ▶ ...
- ▶ **Prevent** undeclared **build time** dependencies.
 - ▶ **Scan** for **runtime** dependencies.
 - ▶ Deploy only **closures** under the **depends-on** relation.

Nix store



Nix store



hello/default.nix

```
{stdenv, fetchurl, perl}:

stdenv.mkDerivation {
  name = "hello-2.1.1";
  builder = ./builder.sh;
  src = fetchurl {
    url =
      ftp://ftp.gnu.org/pub/gnu/hello/hello-2.1.1.tar.gz;
    md5 = "70c9ccf9fac07f762c24f2df2290784d";
  };
  inherit perl;
}
```

hello/default.nix

```
{stdenv, fetchurl, perl}:
```

Function arguments

```
stdenv.mkDerivation {  
  name = "hello-2.1.1";  
  builder = ./builder.sh;  
  src = fetchurl {  
    url =  
      ftp://ftp.gnu.org/pub/gnu/hello/hello-2.1.1.tar.gz;  
    md5 = "70c9ccf9fac07f762c24f2df2290784d";  
  };  
  inherit perl;  
}
```

hello/default.nix

```
{stdenv, fetchurl, perl}:
```

Function arguments

```
stdenv.mkDerivation {  
  name = "hello-2.1.1";  
  builder = ./builder.sh;  
  src = fetchurl  
    url =  
      ftp://ftp.gnu.org/pub/gnu/hello/hello-2.1.1.tar.gz;  
  md5 = "70c9ccf9fac07f762c24f2df2290784d";  
};  
inherit perl;  
}
```

Build attributes

hello/builder.sh

```
. $stdenv/setup

PATH=$perl/bin:$PATH

tar xvfz $src
cd hello-*
./configure --prefix=$out
make
make install
```

hello/builder.sh

```
. $stdenv/setup  
  
PATH=$perl/bin:$PATH  
  
tar xvfz $src  
cd hello-*  
./configure --prefix=$out  
make  
make install
```

Environment initially empty; prevents undeclared dependencies

system/all-packages-generic.nix

```
hello = (import ../applications/misc/hello/ex-1) {  
    inherit fetchurl stdenv perl;  
};  
  
perl = (import ../development/interpreters/perl) {  
    inherit fetchurl stdenv;  
};  
  
fetchurl = (import ../build-support/fetchurl) {  
    inherit stdenv; ...  
};  
  
stdenv = ...;
```


system/all-packages-generic.nix

```
hello = (import ../applications/misc/hello/ex-1) {  
    inherit fetchurl stdenv perl;  
};  
  
perl ← (import ../development/interpreters/perl) {  
    inherit fetchurl stdenv;  
};  
  
fetchurl = (import ../build-support/fetchurl) {  
    inherit stdenv; ...  
};  
  
stdenv = ...;
```

Back to service deployment

We're just going to build service configuration files and control scripts as software components, i.e., as immutable objects in the Nix store.

services/svn.nix

```
{ stdenv, apacheHttpd, subversion }:  
stdenv.mkDerivation {  
  name      = "svn-service";  
  builder   = ./builder.sh; # Build script.  
  control   = ./control.in; # Control script template.  
  conf      = ./httpd.conf.in; # Apache configuration template.  
  inherit  apacheHttpd subversion;  
}
```

Back to service deployment

services/httpd.conf.in

```
...  
LoadModule dav_svn_module  
    @subversion@/modules/mod_dav_svn.so  
...
```

Back to service deployment

services/control.in

```
#! @shell@/bin/sh
...
@apacheHttpd@/sbin/apachectl -k start
    -f @out@/httpd.conf
...
```

services/builder.sh

The builder just replaces **@subversion@**, **@apacheHttpd@**, etc., with the actual paths of the components in the Nix store (passed as arguments to the function).

Using the service

Building and starting

```
# upgrade-server svn ./svn.nix
```

Upgrading and restarting


Idem

Rollback

```
# /nix/var/nix/profiles/svn-service/bin/control stop  
# nix-env -f /nix/var/nix/profiles/svn-service --rollback  
# /nix/var/nix/profiles/svn-service/bin/control start
```

- ▶ The previous service is *monolithic*: the Apache instance provides only the Subversion service.
- ▶ In general it should be possible to run multiple services (“subservices”) within a web server, database server, etc.
- ▶ Therefore the previous service should be split into:
 - ▶ An Apache service
 - ▶ A list of *subservices* that plug into Apache

Example



Example — combined Subversion / TWiki server

```
subversionService = import ../subversion-service {  
    httpPort = 80;  
    reposDir = "/data/subversion"; ...  
};  
twikiService = import ../twiki-service {  
    twikisDir = "/data/twiki"; ...  
};  
webServer = import ../apache-httpd {  
    inherit (pkgs) stdenv apacheHttpd;  
    hostName = "svn.cs.uu.nl";  
    httpPort = 80;  
    subServices = [subversionService twikiService];  
};
```

Example — combined Subversion / TWiki server

```
subversionService = import ../subversion-service {  
    httpPort = 80; # Oops!  
    reposDir = "/data/subversion"; ...  
};  
twikiService = import ../twiki-service {  
    twikisDir = "/data/twiki"; ...  
};  
webServer = import ../apache-httpd {  
    inherit (pkgs) stdenv apacheHttpd;  
    hostName = "svn.cs.uu.nl";  
    httpPort = 80; # Oops!  
    subServices = [subversionService twikiService];  
};
```

Cross-cutting configuration

- ▶ The Nix expression language is functional, which makes it easy to define cross-cutting configuration choices once and propagate them to their realisation sites.
- ▶ This also makes it easy to express variability in configurations.
 - ▶ E.g., whether to build a test or production server
 - ▶ Due to hashing any change to the configuration will result in the resulting components stored in a different location in the Nix store

Example

```
{productionServer}: # Variation point
let {
  port = if productionServer then 80 else 8080;
  webServer = import ./apache-httpd {
    inherit (pkgs) stdenv apacheHttpd;
    hostName = "svn.cs.uu.nl";
    httpPort = port;
    subServices = [subversionService twikiService];
  };
  subversionService = import ./subversion-service {
    httpPort = port;
    reposDir = "/data/subversion"; ...
  };
  twikiService = import ./twiki-service {
    twikisDir = "/data/twiki"; ...
  };
}
```

- ▶ Services frequently consist of subservices running on different machines / platforms

- ▶ Nix already supports multi-platform distributed builds, semi-transparently:

```
derivation {  
  name = "foo";  
  builder = ./builder.sh;  
  system = "i686-linux"; ... }
```

Attribute system denotes platform for component build action; if machine is not **i686-linux**, the build will be forwarded to a machine of the right type.

- ▶ *Starting* and *stopping* is done by a *service runner* component that remotely starts/stops subservices on the machines identified by their **host** attributes.

Example

PostgreSQL server on FreeBSD

Build a Postgres server on FreeBSD.

```
postgresService = import ./postgresql {  
    inherit (pkgsFreeBSD) stdenv postgresql;  
    host = "losser.labs.cs.uu.nl"; # Machine to run on.  
    dataDir = "/var/postgres/jira-data";  
  
    subServices = [jiraService];  
    allowedHosts = [jettyService.host]; # Access control.  
};
```


Example (cont'd)

Jetty container on Linux

Build a Jetty container on Linux.

```
jettyService = import ./jetty {  
  inherit (pkgsLinux) stdenv jetty j2re;  
  host = "itchy.labs.cs.uu.nl"; # Machine to run on.
```

Include the JIRA web application at URI path.

```
  subServices = [ { path = "/jira"; war = jiraService; } ];  
};
```

Build a JIRA service.

```
jiraService = import ./jira/server-pkgs/jira/jira-war.nix {  
  inherit (pkgsLinux) stdenv fetchurl ant postgresql_jdbc;  
  databaseHost = postgresService.host; # Database to use.  
};
```

Example (cont'd)

Service runner on Linux

```
# Compose the two services.  
serviceRunner = import ./runner {  
    inherit (pkgsLinux) stdenv substituter;  
    services = [postgresService jettyService];  
};
```

- ▶ Package management tools (e.g., RPM)
 - ▶ Don't do service configuration
- ▶ Build managers
 - ▶ Make: doesn't deal with variability very well
 - ▶ Better build managers: Odin, Vesta; these have better variability support
- ▶ Cfgengine (Burgess 1995)
 - ▶ Specification of *destructive* action to be performed to realise a desired target state
 - ▶ E.g., add lines *X* to configuration file *Y*
 - ▶ Destructive model and lack of abstraction over paths makes it hard to support multiple instances of a service
 - ▶ Disconnect from software deployment

Conclusion

- ▶ Nix's properties for software deployment carry over to service deployment
- ▶ I.e., full dependencies (\Rightarrow reproducibility), automatic builds, side-by-side deployment of variants, rollbacks
- ▶ Nix expression language is good for dealing with cross-cutting configuration choices
- ▶ Approach extends to distributed deployment

Links

- ▶ Homepage: <http://www.cs.uu.nl/groups/ST/Trace/Nix>
- ▶ Installation instructions for the Subversion service:
<https://svn.cs.uu.nl:12443/repos/trace/services/trunk/subversion/INSTALL>