

TraCE: Transparent Configuration Environments

Martin Bravenboer Eelco Dolstra Eelco Visser

Institute of Information & Computing Sciences
Utrecht University, The Netherlands

February 4, 2005

Project goal

Improving the configuration of software systems

Team

- ▶ Eelco Visser (principal investigator)
- ▶ Martin Bravenboer (PhD student)
- ▶ Eelco Dolstra (PhD student)
- ▶ Gert Florijn (CIBIT)
- ▶ Doaitse Swierstra (promotor)
- ▶ Merijn de Jonge (was postdoc, now at Philips Research)

Variability in Software Systems

Variability is the ability to select the set of features for a particular instance of a software system.

Configuration of variability is realized through configuration mechanisms.

Dimensions of Software Configuration

Type of variability

- ▶ functionality
- ▶ version
- ▶ platform
- ▶ dependencies

Unit of configuration

- ▶ file
- ▶ package
- ▶ system

Moment of configuration

- ▶ development
- ▶ build
- ▶ distribution
- ▶ installation
- ▶ activation
- ▶ run

Different mechanisms at each point in configuration space

Dimensions of Software Configuration

Type of variability

- ▶ functionality
- ▶ version
- ▶ platform
- ▶ dependencies

Unit of configuration

- ▶ file
- ▶ package
- ▶ system

Moment of configuration

- ▶ development
- ▶ build
- ▶ distribution
- ▶ installation
- ▶ activation
- ▶ run

Different mechanisms at each point in configuration space

Dimensions of Software Configuration

Type of variability

- ▶ functionality
- ▶ version
- ▶ platform
- ▶ dependencies

Unit of configuration

- ▶ file
- ▶ package
- ▶ system

Moment of configuration

- ▶ development
- ▶ build
- ▶ distribution
- ▶ installation
- ▶ activation
- ▶ run

Different mechanisms at each point in configuration space

Dimensions of Software Configuration

Type of variability

- ▶ functionality
- ▶ version
- ▶ platform
- ▶ dependencies

Unit of configuration

- ▶ file
- ▶ package
- ▶ system

Moment of configuration

- ▶ development
- ▶ build
- ▶ distribution
- ▶ installation
- ▶ activation
- ▶ run

Different mechanisms at each point in configuration space


Example Problem: Conditional Compilation

```
#ifdef (HP9000_S800) /* If HP9000_S800 is defined, INT_SIZE */
#define INT_SIZE 32 /* is defined to be 32 (bits). */
#elif defined (HPVECTRA) && defined (SMALL_MODEL)
#define INT_SIZE 16 /* Otherwise, if HPVECTRA and */
#endif /* SMALL_MODEL are defined,INT_SIZE is */

#ifdef DEBUG /* If DEBUG is defined, display the */
printf("table element : \n"); /* table elements. */
for (i=0; i < MAX_TABLE_SIZE; ++i)
printf("%d %f\n", i, table[i]);
#endif
```

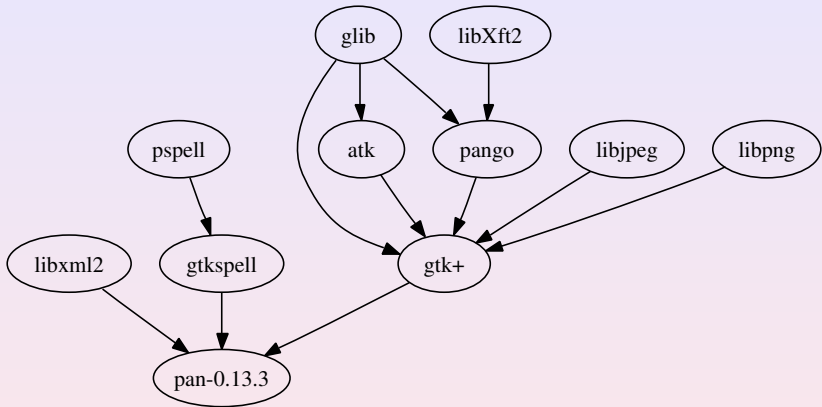
configuration with unhygienic lexical program transformations

Example Problem: Package Dependencies



incomplete dependencies lead to unsafe deployment

Example Problem: Package Dependencies



incomplete dependencies lead to unsafe deployment

Example Problem: Release Management

Distribution



Source distribution

- [strategoxt-0.13.tar.gz](#) (6864337 bytes; MD5 hash: 783bea5d5ebc0604e7ecf5bfb8f7f7b1)



RPM for Red Hat 9.0

- [strategoxt-0.13-1.i386.rpm](#) (18463282 bytes; MD5 hash: 7f0359c78759cc51f864d0961d7f1b57)
- [strategoxt-0.13-1.src.rpm](#) (6818930 bytes; MD5 hash: f762d43367485d97631a62a7266c81dd)

This RPM requires that the following packages are also installed:

- [aterm-2.3.1-1.i386.rpm](#)
- [sdf2-bundle-2.3-1.i386.rpm](#)



RPM for Fedora Core 2

- [strategoxt-0.13-1.i386.rpm](#) (18387992 bytes; MD5 hash: ad714c3594074eeb45f6538f75c4989e)
- [strategoxt-0.13-1.src.rpm](#) (6818933 bytes; MD5 hash: 6a63e79eb94783c01ea0f63e48559538)

This RPM requires that the following packages are also installed:

- [aterm-2.3.1-1.i386.rpm](#)
- [sdf2-bundle-2.3-1.i386.rpm](#)



RPM for Fedora Core 3

- [strategoxt-0.13-1.i386.rpm](#) (18311880 bytes; MD5 hash: 5750c4092d055fb1e846813459885400)
- [strategoxt-0.13-1.src.rpm](#) (6818937 bytes; MD5 hash: df556ab36e20fb3be530748a18f21ebf)

manually releasing many packages often does not scale

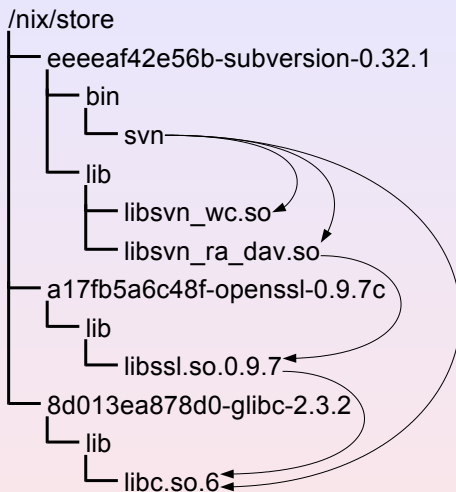
Transparent configuration of software systems
at all levels of granularity and
at all moments on the development/deployment timeline.

Example Solution: User-Interface Configuration

```
public void createLayout(){
    this {
        content = panel {
            border = line border borggreen
            layout = border layout {
                center = label {
                    text = "Welcome"
                    border = raised etched border
                    doublebuffered = true
                }
                south = new JButton("Ok")
            }
        };
    this.pack();
    this.setVisible(true);
}
```

high-level domain-specific configuration supported by hygienic transformations on code structure

Example Solution: Dependency Closures



Safe deployment by computing complete closure of dependency relation

Example Solution: Automatic Release Management

Distribution



Source distribution

- [strategoxt-0.13.tar.gz](#) (6864337 bytes; MD5 hash: 783bea5d5ebc0604e7ecf5bfb8f7f7b1)



RPM for Red Hat 9.0

- [strategoxt-0.13-1.i386.rpm](#) (18463282 bytes; MD5 hash: 7f0359c78759cc51f864d0961d7f1b57)
- [strategoxt-0.13-1.src.rpm](#) (6818930 bytes; MD5 hash: f762d43367485d97631a62a7266c81dd)

This RPM requires that the following packages are also installed:

- [aterm-2.3.1-1.i386.rpm](#)
- [sdf2-bundle-2.3-1.i386.rpm](#)



RPM for Fedora Core 2

- [strategoxt-0.13-1.i386.rpm](#) (18387992 bytes; MD5 hash: ad714c3594074eeb45f6538f75c4989e)
- [strategoxt-0.13-1.src.rpm](#) (6818933 bytes; MD5 hash: 6a63e79eb94783c01ea0f63e48559538)

This RPM requires that the following packages are also installed:

- [aterm-2.3.1-1.i386.rpm](#)
- [sdf2-bundle-2.3-1.i386.rpm](#)



RPM for Fedora Core 3

- [strategoxt-0.13-1.i386.rpm](#) (18311880 bytes; MD5 hash: 5750c4092d055fb1e846813459885400)
- [strategoxt-0.13-1.src.rpm](#) (6818937 bytes; MD5 hash: df556ab36e20fb3be530748a18f21ebf)

completely automatic creation of source and binary distributions

Software deployment (Nix)

SCM'03 Integrating Software Construction and Software Deployment

ICSE'04 Imposing a Memory Management Discipline on Software Deployment

LISA'04 Nix: A Safe and Policy-Free System for Software Deployment.

Build-level composition (AutoBundle/Koala)

ICSR'04 Decoupling Source Trees into Build-Level Components

Code-Level configuration (Stratego/XT)

OOPSLA'04 Concrete Syntax for Objects

Dynamic component composition (XTC)

Thesis'05 Transformation Tool Composition

What do we have to offer?

Ready for showtime

The *Nix Deployment System*

- ▶ ... nice features ...
- ▶ web service configuration and deployment
- ▶ continuous integration
- ▶ automatic release management

Under research

- ▶ Source-level configuration
- ▶ Configuration of compositions

TraCE Workshop

- ▶ Overview of Nix
- ▶ Opportunities for further research and collaboration

Part II

TraCE Workshop

- ▶ Overview of TraCE research topics (past, present and future) with interactive discussion.
 - ▶ Nix
 - ▶ Code-level configuration
- ▶ Questions:
 - ▶ Relevance for industry
 - ▶ Other research questions?
 - ▶ ...

Project goal

Improving the configuration of software systems

Team

- ▶ Eelco Visser (principal investigator)
- ▶ Martin Bravenboer (PhD student)
- ▶ Eelco Dolstra (PhD student)
- ▶ Gert Florijn (CIBIT)
- ▶ Doaitse Swierstra (promotor)
- ▶ Merijn de Jonge (was postdoc, now at Philips Research)

Software deployment (Nix)

SCM'03 Integrating Software Construction and Software Deployment

ICSE'04 Imposing a Memory Management Discipline on Software Deployment

LISA'04 Nix: A Safe and Policy-Free System for Software Deployment.

Build-level composition (AutoBundle/Koala)

ICSR'04 Decoupling Source Trees into Build-Level Components

Code-Level configuration (Stratego/XT)

OOPSLA'04 Concrete Syntax for Objects

Dynamic component composition (XTC)

Thesis'05 Transformation Tool Composition

Problems

Poor and different techniques, e.g.

- ▶ Lexical pre-processing
- ▶ Conditional includes
- ▶ Source file selection
- ▶ Code generation
- ▶ Flexible, but difficult to use and select

Goals

- ▶ Unify configuration mechanisms in source code
- ▶ Improve mapping of configuration from environment to source code

Source-level configuration (2)

Solutions

- ▶ Abstract over concrete configuration mechanisms
 - ▶ Generate concrete use
 - ▶ Select concrete mechanism
- ▶ Separation of configuration issues
 - ▶ Domain-specific languages and generators
 - ▶ Avoid mixing real code and configuration issues: separate meta-level interfaces


Techniques

- ▶ Program manipulation and generation
- ▶ Domain-specific language design, implementation, and embedding

The Nix Deployment System

- ▶ Software deployment: the art of **transferring software** (packages) from one machine to another (and managing it).
- ▶ The hard part: packages should **work the same** on the target machine.
 - ▶ “DLL hell”
 - ▶ “Dependency hell”
 - ▶ Labour-intensive

So why is this hard?



- ▶ Central idea: store all packages in isolation.
- ▶ Unique paths:

```
/nix/store/605332199533e73b...-gtk+-2.2.4
```

which is an MD5 hash of **all** inputs used to build the package:

- ▶ Libraries
 - ▶ Compilers
 - ▶ Build scripts
 - ▶ Build parameters
 - ▶ System type
 - ▶ ...
- ▶ **Prevent** undeclared **build time** dependencies.
 - ▶ **Scan** for **runtime** dependencies.
 - ▶ Deploy only **closures** under the **depends-on** relation.

Advantages

- ▶ Safe deployment
- ▶ Effective composition language
- ▶ User environments
- ▶ Less packaging effort
- ▶ Build farms / automated release management
- ▶ Efficient upgrade deployment
- ▶ Policy-freeness
- ▶ Support for service deployment

- ▶ Hashing gives variability support for free, so no overwriting of dependencies
- ▶ Full dependency graph
- ▶ Complete deployment
- ▶ Side-by-side versioning
- ▶ Automatic garbage collection of unused components

Composition language

- *Nix expressions*: simple functional language to describe components and compositions; easy to express variants

```
bittorrent = (import ../tools/networking/bittorrent) {  
  inherit fetchurl stdenv wxGTK;  
};  
  
wxGTK = (import ../development/libraries/wxGTK) {  
  inherit fetchurl stdenv pkgconfig;  
  gtk = gtkLibs22.gtk;  
};  
  
firefox = (import ../applications/browsers/firefox) {  
  inherit fetchurl stdenv pkgconfig perl zip libIDL libXi;  
  gtk = gtkLibs24.gtk;  
};
```

- ▶ *User environments* are components of symlinks to activated components
- ▶ Can be per-user/process/service/...
- ▶ Allow atomic upgrading / downgrading (rollbacks)

Less packaging effort

- ▶ Nix expressions describe a source deployment model
- ▶ No need to explicitly to do binary packaging due to *transparent binary deployment*: binaries can be *cached* in a shared repository

- ▶ Nix is a good basis for a build farm because:
 - ▶ Expression language ideal for describing build tasks
 - ▶ Expression language makes it easy to describe variant compositions
 - ▶ Nix manages the dependencies
 - ▶ Supports for distributed builds
 - ▶ Hashing scheme + complete dependencies allow builds to be reproduced reliably
 - ▶ Efficiency: due only rebuild things that have changed
- ▶ Useful for:
 - ▶ Continuous integration testing
 - ▶ Portability testing
 - ▶ *Automated release management* — successful builds are releases

- ▶ Successful builds are automatically *released* as stable or unstable releases.
- ▶ Releases can be automatically *pushed to* / *pulled by* clients.

KoalaCompiler release koala-compiler-0.1pre8399

This is a **bad** release: one or more of its build steps failed. See [below](#) for details. This release should not be used for production purposes.

This page provides release **koala-compiler-0.1pre8399** of KoalaCompiler. It was generated automatically on 2004-12-22 19:59:06 UTC from revision 8399 of the path [trunk/koala-compiler](#) of its Subversion repository (the [XML record of the build job](#) is available).

Distribution



Source distribution

- [koala-compiler-0.1pre8399.tar.gz](#) (1389572 bytes; MD5 hash: e99278ec393b979ad06561e9cd626c80)



RPM for Red Hat 9.0

- [koala-compiler-0.1pre8399-1.i386.rpm](#) (2282473 bytes; MD5 hash: b9e9094dfcdefc29704a83b8d563b83d)
- [koala-compiler-0.1pre8399-1.src.rpm](#) (1379161 bytes; MD5 hash: 354d4ddba68273c4d73d4669ae7140ea)

This RPM requires that the following packages are also installed:

- [aterm-2.2-1.i386-redhat9.0-linux-gnu.rpm](#)
- [sdf2-bundle-2.2.i386-redhat9.0-linux-gnu.rpm](#)
- [strategoxt-0.13pre8212-1.i386.rpm](#)



SuSE RPM for SuSE 9.0

- [koala-compiler-0.1pre8399-1.i586.rpm](#) (2334444 bytes; MD5 hash: 58dd1fc0341aede672e90cf2e0c8c84a)
- [koala-compiler-0.1pre8399-1.src.rpm](#) (1379158 bytes; MD5 hash: b6a9c0e22744bd03eab81196eb78cd77)

Release management

- ▶ Successful builds are automatically *released* as stable or unstable releases.
- ▶ Releases can be automatically *pushed to* / *pulled by* clients.

```
making check in koala-0
+ building check-recursive
Making check in koala-stc
- building check-recursive
- make[2]: Entering directory `/tmp/nix-22047-22/koala-compiler-0.1pre8399/demos/koala-stc'
- Making check in sgldr-flat-bundle
- building check-am
- make[3]: Entering directory
  `/tmp/nix-22047-22/koala-compiler-0.1pre8399/demos/koala-stc/sgldr-flat-bundle'
- make check-TESTS
- building test
- make[4]: Entering directory
  `/tmp/nix-22047-22/koala-compiler-0.1pre8399/demos/koala-stc/sgldr-flat-bundle'
- building check-TESTS
- 3,7d2
- < ./test/sgldr_bundle/AUTHORS
- < ./test/sgldr_bundle/ChangeLog
- < ./test/sgldr_bundle/Makefile.am
- < ./test/sgldr_bundle/NEWS
- < ./test/sgldr_bundle/README
- 10a6,7
- > ./test/sgldr_bundle/AUTHORS
- > ./test/sgldr_bundle/ChangeLog
- 15a13,14
- > ./test/sgldr_bundle/Makefile.am
- > ./test/sgldr_bundle/NEWS
- 17a17
- > ./test/sgldr_bundle/README
- FAIL: ../../demos/koala-stc/test-koala
- =====
- 1 of 1 tests failed
- Please report to mdejonge@cs.uu.nl
- =====
- make[4]: *** [check-TESTS] Error 1
- make[4]: Leaving directory
  `/tmp/nix-22047-22/koala-compiler-0.1pre8399/demos/koala-stc/sgldr-flat-bundle'
- make[3]: *** [check-am] Error 2
- make[3]: Leaving directory
```

- ▶ Successful builds are automatically *released* as stable or unstable releases.
- ▶ Releases can be automatically *pushed to* / *pulled by* clients.

KoalaCompiler release koala-compiler-0.1

This page provides release **koala-compiler-0.1** of KoalaCompiler. It was generated automatically on 2004-12-22 21:14:12 UTC from revision 8401 of the path `/trunk/koala-compiler` of its Subversion repository (the [XML record of the build job](#) is available).

Distribution



Source distribution

- [koala-compiler-0.1.tar.gz](#) (1388530 bytes; MD5 hash: 06fc55524399a17e95705c5ddf4b406f)



RPM for Red Hat 9.0

- [koala-compiler-0.1-1.i386.rpm](#) (2282430 bytes; MD5 hash: 475e1e609d5e3e1b29fe97bb0e60213b)
- [koala-compiler-0.1-1.src.rpm](#) (1379515 bytes; MD5 hash: 2a0fbdc6ae1a14af6f1db012ade40077)

This RPM requires that the following packages are also installed:

- [aterm-2.2-1.i386-redhat9.0-linux-gnu.rpm](#)
- [sdf2-bundle-2.2.i386-redhat9.0-linux-gnu.rpm](#)
- [strategox-0.13pre8212-1.i386.rpm](#)



SuSE RPM for SuSE 9.0

- [koala-compiler-0.1-1.i586.rpm](#) (2334410 bytes; MD5 hash: 517555f32098effc53676132cb3d4490)
- [koala-compiler-0.1-1.src.rpm](#) (1379517 bytes; MD5 hash: d3171e0fcfe54f8321dbfa0c7ced468d)

This RPM requires that the following packages are also installed:

- [aterm-2.2-1.i386-redhat9.0-linux-gnu.rpm](#)
- [sdf2-bundle-2.2.i386-redhat9.0-linux-gnu.rpm](#)


Efficient upgrade (patch) deployment

- ▶ New versions of components can be deployed as *patches* from older versions
- ▶ Completely transparent to users; downloader selects shortest sequence of full or patch downloads
- ▶ Also transparent to packagers; integrated into release management

- ▶ Easy to define new deployment policies, e.g.,
 - ▶ Push/pull models
 - ▶ Channels
 - ▶ Manual
 - ▶ Multi-level
 - ▶ Whether to do source-only, binary-only, source/binary deployment
 - ▶ What variants to pre-build
 - ▶ ...

Service deployment

- ▶ Deploying a service is (almost) the same as deploying software.
- ▶ Example: Subversion server at **svn.cs.uu.nl**.



- ▶ Stable implementation.
- ▶ Open source.
- ▶ Available at <http://www.cs.uu.nl/groups/ST/Trace/Nix>
- ▶ Also: Nix Packages collection; large set of common Linux components
- ▶ Documentation:
 - ▶ Manual
 - ▶ Several papers (ICSE-2004, LISA-2004, submitted paper on patch deployment)

- ▶ Security aspects
 - ▶ Multi-user stores
 - ▶ Ensuring security fix deployment
- ▶ “Low-level” build management
- ▶ Deployment for distributed systems; mass deployment (incl. grids)
- ▶ Configuration selection; e.g., automatically finding configurations meeting constraints